

É

C

A

R

T

Endianness and Network Byte Order

Big Endian vs Little Endian

- When sending a word (2 or 4 bytes), the reading/writing order is important
- Endianness refers to the order of the bytes, comprising a digital word, in computer memory. Definitions from Wikipedia
- Big Endian: the most significant byte of a word is stored in a particular memory address, and subsequent bytes are stored in the following higher memory addresses
- Little Endian: the least significant byte of a word is stored in a particular memory address, and subsequent bytes are stored in the following higher memory addresses
- Network byte order is Big-Endian

Communication without Network

Byte Order

Short Integer = $255_{10} = 00\ FF_{16}$

Memory (X, X+1) = (FF,00)

Write 1st Byte = FF

Write 2nd Byte = 00



Little-Endian

00 FF

Receives 1st Byte = FF

Receives 2nd Byte = 00

Memory (X,X+1) = (FF,00)

Unsigned short Integer = $65280_{10} = FF\ 00_{16}$



Big-Endian

Network Byte Order formatting in C

- 2 bytes words are
 - written in network byte order with the `htons()`. If the device is BE, no actions are taken, otherwise, bytes are flipped
 - Translated to the device architecture with `ntohs()`. If the device is BE, no actions are taken, otherwise, bytes are flipped
- 4 bytes words are
 - written in network byte order with the `htonl()`. Same observations as above.
 - Translated to the device architecture with `ntohl()`. Same observations as above.
- 1-byte data is not impacted by the network byte order

Network Byte Order formatting in Python

- `struct.pack()` to pack binary data, convert to the network byte order and send it by the socket
- `struct.unpack()` to receive bytes

Is my computer big or little endian?

```
from math import ceil
from struct import pack
```

```
def show_bytes(data):
```

```
    i = 0
    for b in data:
        print("Byte %d has %02x" % (i,b))
        i=i+1
```

```
var = int("16909060",10)
numBytes = ceil(var.bit_length()/8.0)
print("Var has %08x" % (var))
```

```
i = 0
while numBytes > i:
    print("Byte %d has %02x" % (i,(var>>(i*8) & 0xff)))
    i=i+1
```

```
print("")
show_bytes(pack(">I",var)); # Big-Endian
```

```
print("")
show_bytes(pack("!I",var)); # Network Byte Order
```

Var has 01020304

Byte 0 has 04

Byte 1 has 03

Byte 2 has 02

Byte 3 has 01

Byte 0 has 01

Byte 1 has 02

Byte 2 has 03

Byte 3 has 04

Byte 0 has 01

Byte 1 has 02

Byte 2 has 03

Byte 3 has 04

Sockets Programming with Python

Architecture & Réseaux

Dino Lopez Pacheco

<http://www.i3s.unice.fr/~lopezpac/>

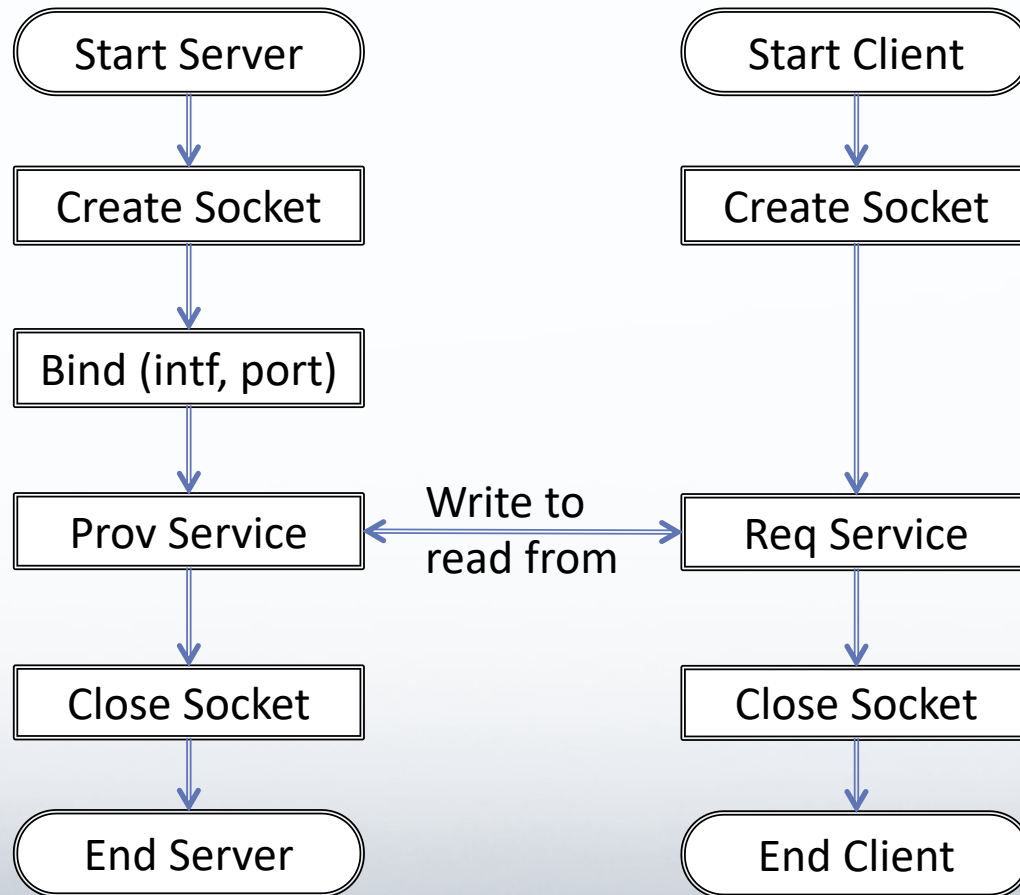
Socket definition

- The sockets are the end points for the communication between 2 processes (Inter-Process Communications – IPC)
 - Local IPC
 - `$ ls ~ | grep "^d" | wc -l`
 - Remote IPC
 - BSD sockets
- Berkeley Sockets (BSD sockets) is a library to allow the programming of Internet Sockets
- BSD sockets evolved and make part now of the POSIX standard
- The Python `socket` module provides access to the BSD Socket interface

Communication modes in IP networks

- Connectionless mode
 - Uses UDP
 - Unreliable
 - Datagram Sockets
- Connection oriented
 - Uses TCP
 - Reliable
 - Stream sockets

Connectionless mode



Writing your code

1. Create your socket – the `socket()` function

- `Socket()` returns a socket object which implements the BSD Socket system calls
- Definition `socket.socket([family[, type[, proto]]])`
 - Address family: by default, `AF_INET`
 - Socket type: by default, `SOCK_STREAM`
 - Protocol number: 0 (zero) frequently

Socket domains and Types

- Several domains
 - *AF_INET: Socket in the IPv4 domain*
 - AF_INET6: Socket in the IPv6 domain
 - AF_BLUETOOTH: Socket in the Bluetooth domain (needs python-bluez)
 - ...
- 3 socket types available for the AF_INET domain
 - *SOCK_STREAM: Connection-oriented communication - TCP*
 - *SOCK_DGRAM: connectionless communication – UDP*
 - SOCK_RAW: custom construction of headers
- All these constants are available in the socket class

Binding the socket with bind()

- bind() to bind an IP address and define a listening port in a newly created socket
- According to the Python doc - `socket.bind(address)`
 - Note that the format of address depend on the address family used to create your socket
 - For AF_INET, the address is a tuple (host,port), where
 - host is a string representing the IP address or the canonical name of a given interface: “mycomp.test.com”, “192.168.0.12”
 - To leave the kernel to take any available interface, use **None** for host
 - Port is an integer
- Bind is mandatory at the server side, but optional at the client side

Sending/Receiving data – connectionless mode

- To send data in a non connected socket - `socket.sendto(bytes, address)`
 - `bytes` represents the message to be sent
 - `address` is the tuple representing the remote host and port
 - It returns the number of bytes which were sent
- To receive the data - `socket.recvfrom(bufsize[, flags])`
 - It returns a pair (`bytes`, `address`), where `bytes` represents the received data and `address` is the address of the remote peer

Closing a Socket

- To close the socket, you can either call the `socket.close()` or the `socket.shutdown(how)` method
- After `close()`, any operation on the socket object will fail
- `shutdown()` allows a finer control over the socket
 - If `how` is `SHUT_RD`, the reception of data is disabled
 - If `how` is `SHUT_WR`, data transmission is disabled
 - If `how` is `SHUT_RDWR`, the transmission and reception of data are disallowed
 - On some OSs, shutting down a half of the connection can close the opposite half
- *In connected mode, closing a socket triggers the transmission of EOF to the remote peer*

Example of a connectionless communication

Connectionless mode – the integer is transmitted in Network Byte Order

Server

```
1. import socket
2. import struct

3. HOST = '' # any available interf
4. PORT = 5000 # Arbitrary non-priv port
5. s = socket.socket(socket.AF_INET,
    socket.SOCK_DGRAM)
6. s.bind((HOST, PORT))
7. data, (HOST, PORT) = s.recvfrom(1024)

8. ndata = struct.unpack("!h", data)
8. data = ndata[0]+1

9. s.sendto(struct.pack("!h", data), (HOST,
    PORT))
10. s.close()
```

Client

```
1. import socket
2. import struct

3. HOST = '127.0.0.1' # The remote host
4. PORT = 5000 # The remote port
5. s = socket.socket(socket.AF_INET,
    socket.SOCK_DGRAM)

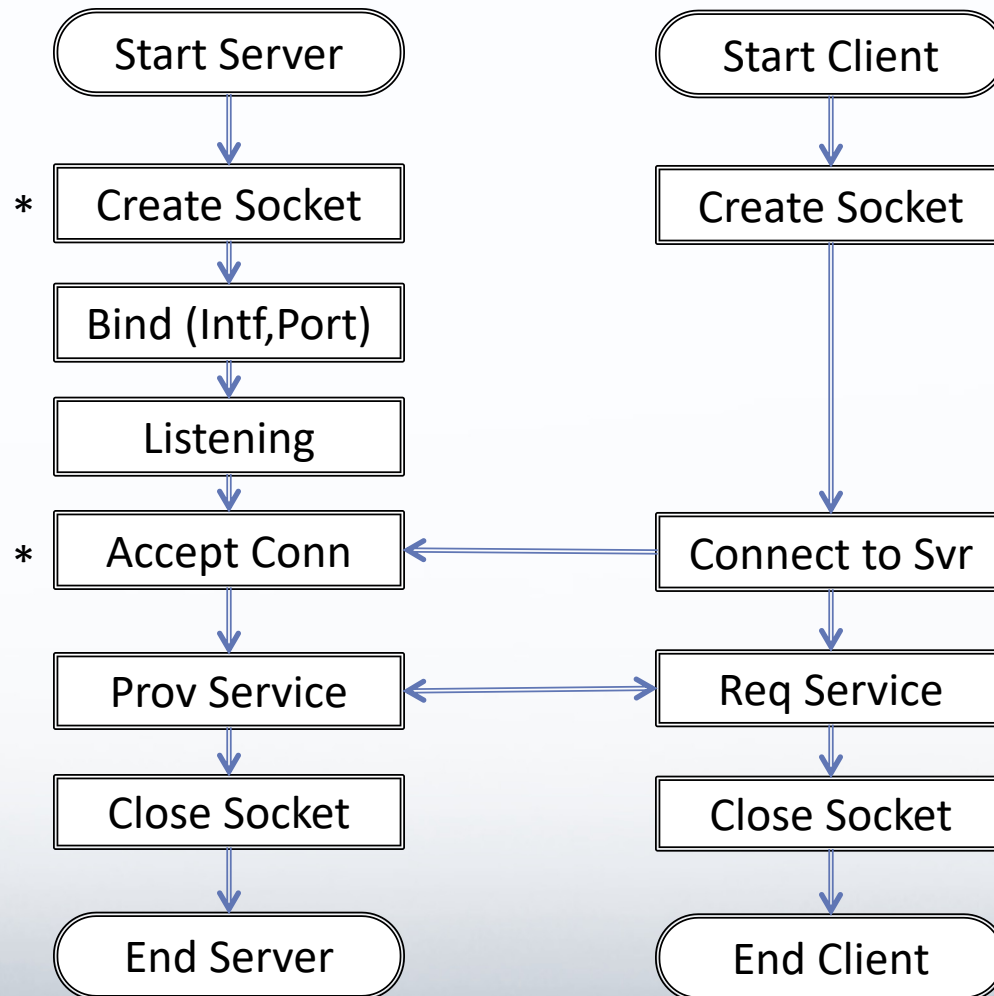
6. val = struct.pack("!h", 1)
7. s.sendto(val, (HOST, PORT))

8. data, (HOST, PORT) = s.recvfrom(1024)
9. print("Received %d" %
    (struct.unpack("!h", data)))

10. s.close()
```

Connection-based communication

The flow chart



Connection request

- In a connection-based communication, after binding the socket, the server must listen for incoming connection.
 - `socket.listen(backlog)`. *backlog* represents the number of incoming connection that can be queued at any time
- After listening for incoming connection, connection requests should be accepted
 - `socket.accept()`. `accept()` returns a pair *conn*, *address*, where
 - *conn* is a new socket object that the server will use to communicate with the remote host
 - *address* is the address of the remote host. The address format depends on the address family type
- The client connect to the server with the `connect()` method.
 - `socket.connect(address)`. *address* is the address of the remote host. The address format depends on the address family type

Sending data

- To send data - `socket.sendall(bytes[, flags])`
 - Send all data unless an error occur
 - It returns “None” on success. Otherwise, an exception is raised and there is no way to know how many data has been sent
 - The optional *flags* argument can be used to execute special sending methods (e.g. out-of-band data)

Receiving data

- To receive data through a connected socket you can use `socket.recv(bufsize[, flags])`
 - `Bufsize` represents the maximum amount of bytes to read
 - `Flags` can be used to perform “special” readings
 - It returns a byte object with the read data

Connection-based communication

Server

```
1. import socket
2. HOST = '' # any available interf
3. PORT = 5000 # Arbitrary non-priv port
4. s = socket.socket(socket.AF_INET,
   socket.SOCK_STREAM)
5. s.bind((HOST, PORT))
6. s.listen(1)
7. conn, addr = s.accept()
8. while 1:
9.     data = conn.recv(1024)
10.    if not data: break
11.    conn.sendall(("Hi!").encode())
12. conn.close()
13. s.close()
```

Client

```
1. import socket
2. HOST = '10.0.0.2' # The
   remote host
3. PORT = 5000 # The
   remote port
4. s =
   socket.socket(socket.AF_INET,
   socket.SOCK_STREAM)
5. s.connect((HOST, PORT))
6. val = "Hello!"
7. s.sendall(val.encode())
8. data = s.recv(1024)
9. s.close()
10. print("Received:
    %s"%(data.decode()))
```

Server styles

- Single Iterating server
 - Only one socket is opened at a time
 - Clients are accepted one after the other
 - Slow service
- Multiprocessing server
 - After accept, the server (fork) creates a subprocess which will provide the service
 - The subprocess is a copy of the parent process
 - The used memory space is doubled
- Multithreading server (POSIX Threads)
 - The same memory space is shared between all the threads
 - Special attention must be taken to avoid race conditions
- Concurrent single server
 - Uses kernel space techniques to simultaneously wait over the whole set of opened socket IDs
 - The main process is wakened up when new data arrives
 - It cannot benefit from multiprocessors

Some thoughts about Multiprocessing

Multiprocessing in Linux / Unix-like systems

- Quick overview of sub-process creation and termination
- One can use the multiprocessing Python package to create and handle sub-processes
 - High-level API
 - Let's play with the OS services to understand the concepts
- `os.fork()` spawns a child process
 - Returns 0 in the child process
 - Returns the PID of the child in the parent process
 - In case of error, an `OSError` exception is raised

Example 1 – No synch between processes

```
1. import os
2. import time

3. def testdelay():
4.     for i in range(0,5):
5.         time.sleep(2)
6.         print("child is
7.             running... %d" %(i))
8.         os._exit(0)

8. pid = os.fork()
9. if pid == 0:
10.     testdelay()

11. print("parent is exiting...")
12. exit(0)
```

parent is exiting...
child is running... 0
child is running... 1
child is running... 2
child is running... 3
child is running... 4

Example 2- waiting for child termination

```
1. import os
2. import time

3. def testdelay():
4.     for i in range(0,5):
5.         time.sleep(2)
6.         print("child is running... %d"
7.             %(i))
8.         os._exit(0)

9. pid = os.fork()
10. if pid == 0:
11.     testdelay()
12. else:
13.     status = os.wait()
14.     print(status)

15. print("parent is exiting...")
16. exit(0)
```

child is running... 0
child is running... 1
child is running... 2
child is running... 3
child is running... 4
(59692, 0)
parent is exiting...

Example 3 – child exits faster than the parent process

```
1. import os
2. import time

3. def testdelay():
4.     for i in range(0,5):
5.         time.sleep(2)
6.         print("parent is
  running... %d" %(i))

7. pid = os.fork()
8. if pid == 0:
9.     print("child exits...")
10.    os._exit(0)
11. else:
12.     testdelay()
13.     status = os.wait()
14.     print(status)

15. print("parent is exiting...")
16. exit(0)
```

- Output
child exits...
parent is running... 0
parent is running... 1
parent is running... 2
parent is running... 3
parent is running... 4
(59692, 0)
parent is exiting...

- Process table status after child finishes but before parent finishes

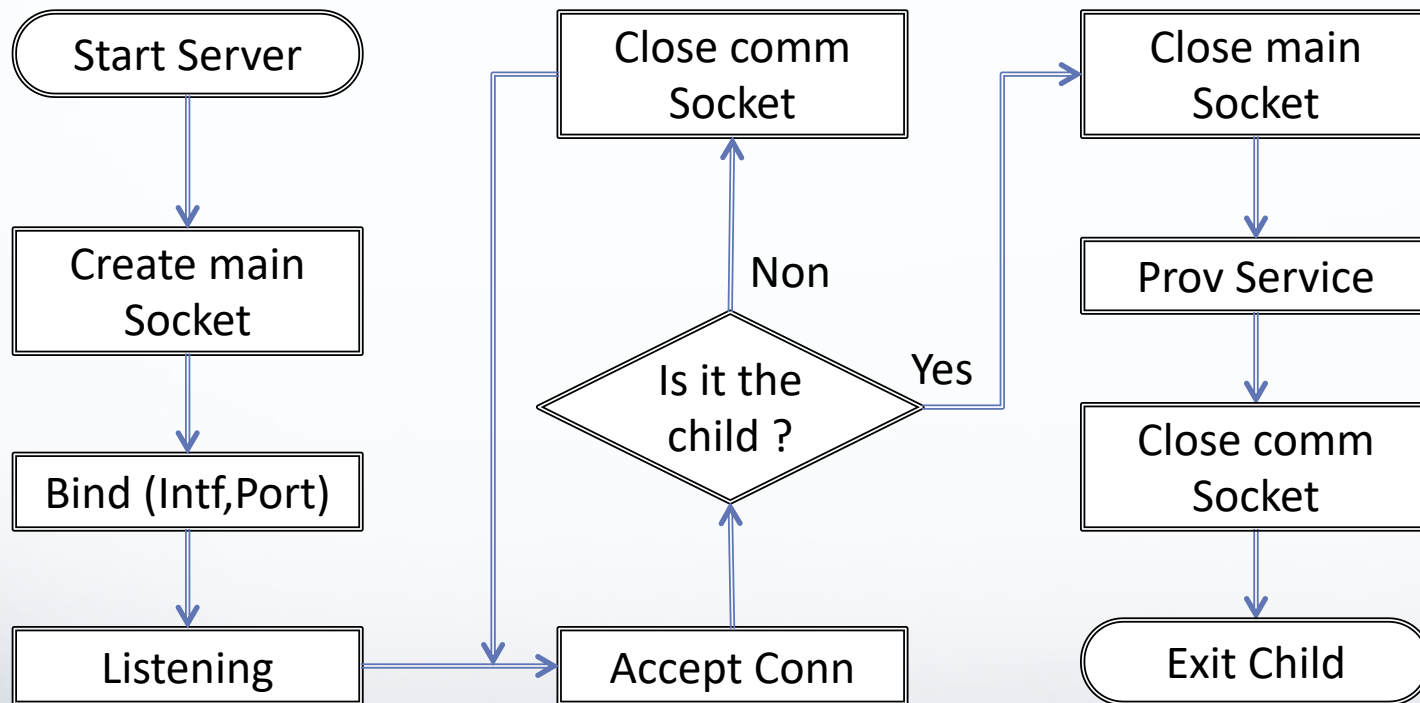
59691	ttys002	0:00.03	python3 test-fork.py
59692	ttys002	0:00.00	[python3] <defunct>

→ Zombie process

SIGCHLD and SIG_IGN

- Upon exit, a child process reports its exit code to its parent
- While the process parent doesn't read the exit status of the child process, this last is keep in the process table
 - Leading to a so-called zombie process
 - Refer to `wait()`, `waitpid()`
- In Linux, Unix-like systems, whenever something interesting happens to a forked off child, the parent process receives a SIGCHLD signal
- By default, SIGCHLD is ignored
- To avoid zombie process, the parent should handle the SIGCHLD signal. Ex.
 - `signal(signal.SIGCHLD, signal.SIG_IGN)`

Multi-process server



Concurrent Single Server

- Event multiplexing/notification
 - Which socket is ready for reading/writing events
- Several ways
 - `select()/poll()`
 - `epoll()` / `kqueue()` → faster than `select/poll`

Concurrent one-thread server

Create svr socket, non blocking

Bind, Listen

Create an epoll object

Add the main socket to epoll - reading mode

Create data collections for “existing connections”, “requests” and “responses”

while true {

wait for events at epoll

for each event { ; // event = fd where the event happens and the event_type

if (fd == svr) {

accept_new_conn ()

} else if (event_type == incoming message) {

handle_reading (fd)

} else if (event_type == outgoing message) {

handle_writing (fd)

}

}

}

The corresponding Python code

```
import socket
import select
```

```
HOST = " " # any available interf
PORT = 5000 # Arbitrary non-priv port
```

```
connections = {}
requests = {}
responses = {}
```

```
e = select.epoll()
s = socket.socket(socket.AF_INET,
                  socket.SOCK_STREAM)
s.setblocking(0)
s.bind((HOST, PORT))
s.listen(1)
```

```
sfd = s.fileno()
e.register(sfd,select.EPOLLIN)
```

```
while 1:
    events = e.poll(1)
    for fileno, event in events:
        if fileno == sfd:
            accept_new_conn()
        elif event & select.EPOLLIN:
            handle_reading(fileno)
        elif event & select.EPOLLOUT:
            handle_writing(fileno)
```

```
s.close()
```


Accept new connections

```
function accept_new_conn() {  
    accept new connection – non  
    blocking  
  
    register the new socket at epoll -  
    reading  
  
    add the new socket at « existing  
    connections »  
  
    create the incoming buffer for  
    requests  
  
    create the outgoing buffer for  
    responses  
  
}
```

```
def accept_new_conn():  
    conn, addr = s.accept()  
    conn.setblocking(0)  
    fd = conn.fileno()  
    e.register(fd, select.EPOLLIN)  
    connections[fd] = conn  
    requests[fd] = None  
    responses[fd] = None
```

Handle input data

```
function handle_reading(fd) {  
    read socket at fd  
    if data read == EOF {  
        unregister fd from epoll  
        close the socket  
        remove fd from existing connections  
        and destroy buffers  
        return  
    }  
    store data at the requests buffer of fd  
    process request  
    store reply at the responses buffer of  
    fd  
    declare an output event for fd  
}
```

```
def handle_reading(fd):  
    data = connections[fd].recv(1024)  
    if not data:  
        e.unregister(fd)  
        connections[fd].close()  
        del connections[fd],  
            requests[fd], responses[fd]  
    return  
    responses[fd] = data;  
    requests[fd] = None  
    e.modify(fd,select.EPOLLOUT)
```

Handle output data

```
function handle_writing(fd) {  
    send the buffer content for fd through the socket  
    clear the output buffer  
    declare fd as ready for reading  
}
```

```
def handle_writing(fd):  
    connections[fd].sendall(responses[fd])  
    responses[fd] = None  
    e.modify(fd,select.EPOLLIN)
```