

Sockets Programming with Python

Réseaux : Programmation & Configuration

Dino Lopez

<http://www.i3s.unice.fr/~lopezpac/>

Socket definition

- The sockets are the end points for the communication between 2 processes (Inter-Process Communications – IPC)
 - Local IPC
 - `$ ls ~ | grep "^d" | wc -l`
 - Remote IPC
 - BSD sockets
- Berkeley Sockets (BSD sockets) is a library to allow the programming of Internet Sockets
- BSD sockets evolved and make part now of the POSIX standard
- The Python `socket` module provides access to the BSD Socket interface

IPC in a nutshell

Receiver (Server)

- Something to read
- Open the recipient
- Define the communication method
- Wait for a call

Sender (Client)

- Something to write
- Identify the destination
- Identify the recipient
- Define the communication method
- Make the call

Client – Server Model

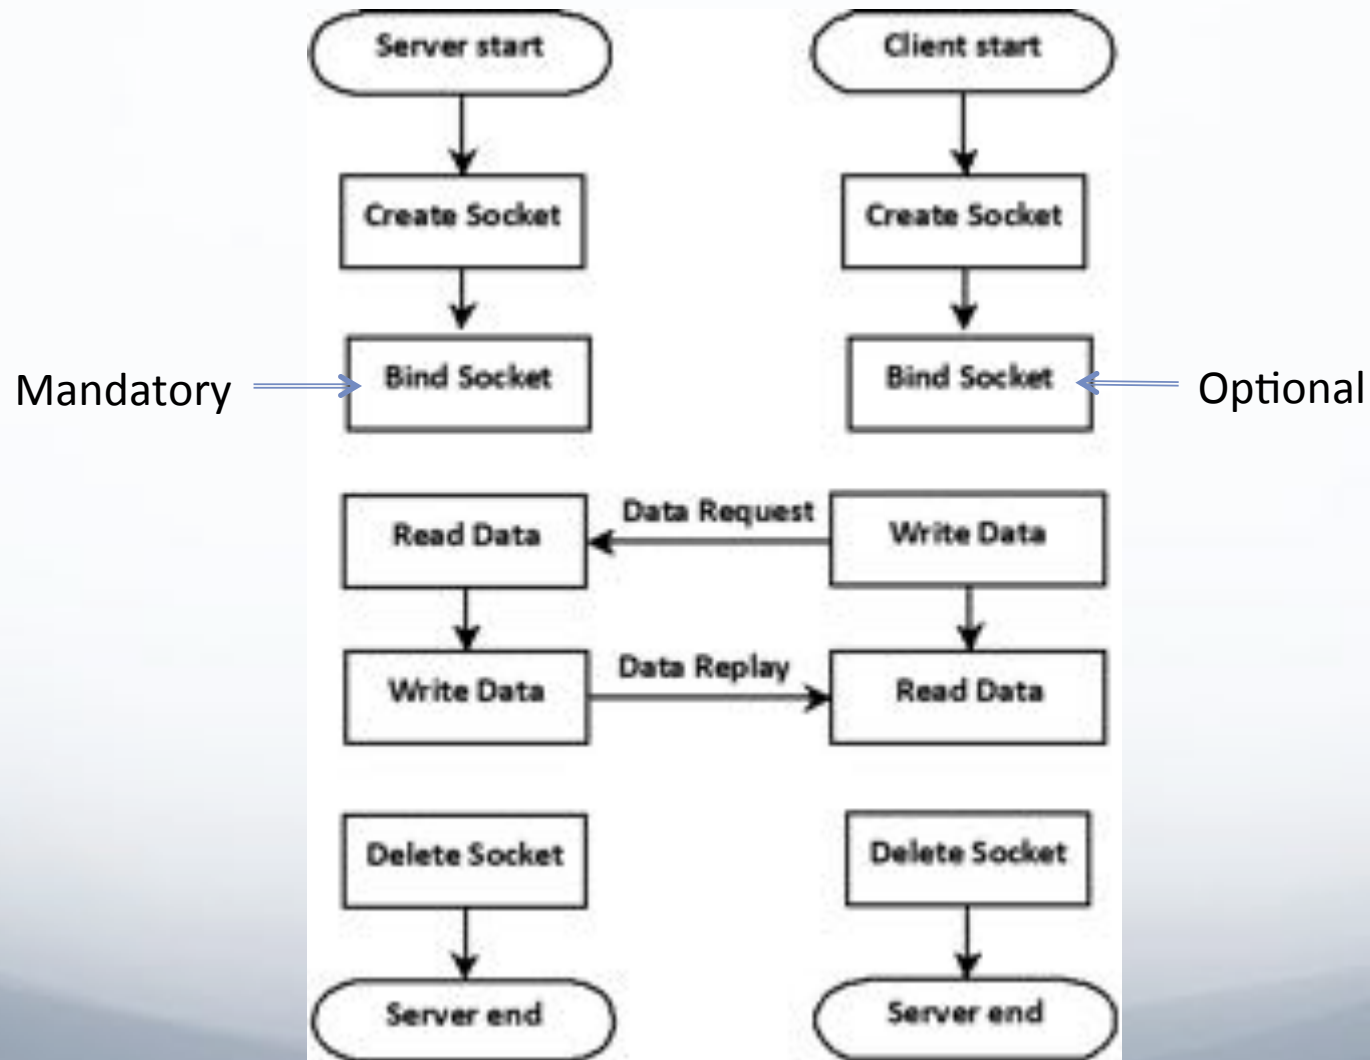
- Server
 - Daemon
- Client
 - Initialize a connection
 - Punctual communications



Communication modes

- Connectionless mode
 - Uses UDP
 - Unreliable
 - Datagram Sockets
- Connection oriented
 - Uses TCP
 - Reliable
 - Monopolize a socket descriptor

Connectionless mode



Writing your code

1. Create your socket – the `socket()` function

- `Socket()` returns a socket object which implements the BSD Socket system calls
- Definition `socket.socket([family[, type[, proto]]])`
 - Address family: by default, `AF_INET`
 - Socket type: by default `SOCK_STREAM`
 - Protocol number: 0 (zero) frequently

Socket domains and Types

- Several domains
 - *AF_INET: Socket in the IPv4 domain*
 - AF_INET6: Socket in the IPv6 domain
 - AF_BLUETOOTH: Socket in the Bluetooth domain (needs python-bluez)
 - ...
- 3 socket types available for the AF_INET domain
 - *SOCK_STREAM: Connection-oriented communication - TCP*
 - *SOCK_DGRAM: connectionless communication – UDP*
 - SOCK_RAW: custom construction of headers
- All these constants are available in the socket class

Binding the socket with bind()

- To bind an IP address and define a listening port in a newly created socket, you will use the bind() method of the socket class
- According to the Python doc, bind() is declared as `socket.bind(address)`
 - Note that the format of address depend on the address family used to create your socket
 - For AF_INET, the address is a tuple (host,port), where
 - host is a string representing the IP address or the canonical name of a given interface: “mycomp.test.com”, “192.168.0.12”
 - To leave the kernel to take any available interface, use **None** for host
 - Port is an integer
- Bind is mandatory at the server side, but optional at the client side

Sending/Receiving data – connectionless mode

- To send data in a non connected socket, you want to use `socket.sendto(string, address)`
 - String represents the message to be sent
 - Address is the tuple representing the remote host and recipient
 - It returns the number of bytes which were actually sent
- And to receive the data, you want to use `socket.recvfrom(bufsize[, flags])` or `socket.recvfrom_into(buffer[, nbytes[, flags]])`
- Regarding `socket.recvfrom(bufsize[, flags])`
 - It returns a pair (string, address), where string represents the received data and address is the address of the remote peer

Closing a Socket

- To close the socket, you can either call the `socket.close()` or the `socket.shutdown(how)` method
- After `close()`, any operation on the socket object will fail
- `shutdown()` allows a finer control over the socket
 - If `how` is `SHUT_RD`, the reception of data is disabled
 - If `how` is `SHUT_WR`, data transmission is disabled
 - If `how` is `SHUT_RDWR`, the transmission and reception of data are disallowed
 - On some OSs, shutting down a half of the connection can close the opposite half
- *In connected mode, closing a socket triggers the transmission of EOF to the remote peer*

Endianess and Network Byte Order

Big Endian vs Little Endian

- Endianness refers to the order of the bytes, comprising a digital word, in computer memory. Definitions from Wikipedia
- Big Endian: the most significant byte of a word is stored in a particular memory address, and subsequent bytes are stored in the following higher memory addresses
- Little Endian: the least significant byte of a word is stored in a particular memory address, and subsequent bytes are stored in the following memory addresses
- Network byte order is Big-Endian

Example in an Intel x86_64 processor

```
from math import ceil
from struct import pack
```

```
def show_bytes(data):
    i = 0
    for b in data:
        print "Byte %d has %02x" % (i,ord(b))
        i=i+1
```

```
var = int("16909060",10)
numBytes = ceil(var.bit_length()/8.0)
print "Var has %08x" % (var)
```

```
i = 0
while numBytes > i:
    print "Byte %d has %02x" % (i,(var>>(i*8) & 0xff))
    i=i+1
```

```
print ""
show_bytes(pack(">I",var)); // Big-Endian
```

```
print ""
show_bytes(pack("!I",var)); // Network Byte Order
```

Var has 01020304
Byte 0 has 04
Byte 1 has 03
Byte 2 has 02
Byte 3 has 01

Byte 0 has 01
Byte 1 has 02
Byte 2 has 03
Byte 3 has 04

Byte 0 has 01
Byte 1 has 02
Byte 2 has 03
Byte 3 has 04

Communication without Network

Byte Order

Short Integer = $255_{10} = 00\text{ FF}_{16}$
Memory (X, X+1) = (FF, 00)
Write 1st Byte = FF
Write 2nd Byte = 00



Little-Endian

00 FF



Big-Endian

Receives 1st Byte = FF
Receives 2nd Byte = 00
Memory (X, X+1) = (FF, 00)
Short Integer = $65280_{10} = \text{FF } 00_{16}$

Network Byte Order formatting

- When sending 2 byte words or 4 byte words, the network byte order must be applied
 - Translated to the device architecture with `ntohl()`. Same observations as above.
- 2 bytes words are
 - written in network byte order with the `htons()`. If the device is BE, no actions are taken, otherwise, bytes are flipped
 - Translated to the device architecture with `ntohs()`. If the device is BE, no actions are taken, otherwise, bytes are flipped
- 4 bytes words are
 - written in network byte order with the `htonl()`. Same observations as above.
- Python is little bit special. `Socket.htonl()`, etc are available. It's preferable to use
 - `struct.pack` to pack binary data, convert to the network byte order and send it by the socket
 - `Struct.unpack` to receive bytes
- When the data has always the right order (e.g. ascii text, file content, etc.) the network byte order does not apply

Example of a connectionless communication

Connectionless mode – the integer is transmitted in Network Byte Order

Server

```
1. import socket
2. import struct

3. HOST = '' # any available interf
4. PORT = 5000 # Arbitrary non-priv port
5. s = socket.socket(socket.AF_INET,
                     socket.SOCK_DGRAM)
6. s.bind((HOST, PORT))
7. data, (HOST, PORT) = s.recvfrom(1024)

8. ndata = struct.unpack("!h", data)
8. data = ndata[0]+1

9. s.sendto(struct.pack("!h", data),
            (HOST, PORT))
10. s.close()
```

Client

```
1. import socket
2. import struct

3. HOST = '127.0.0.1' # The remote host
4. PORT = 5000 # The remote port
5. s = socket.socket(socket.AF_INET,
                     socket.SOCK_DGRAM)

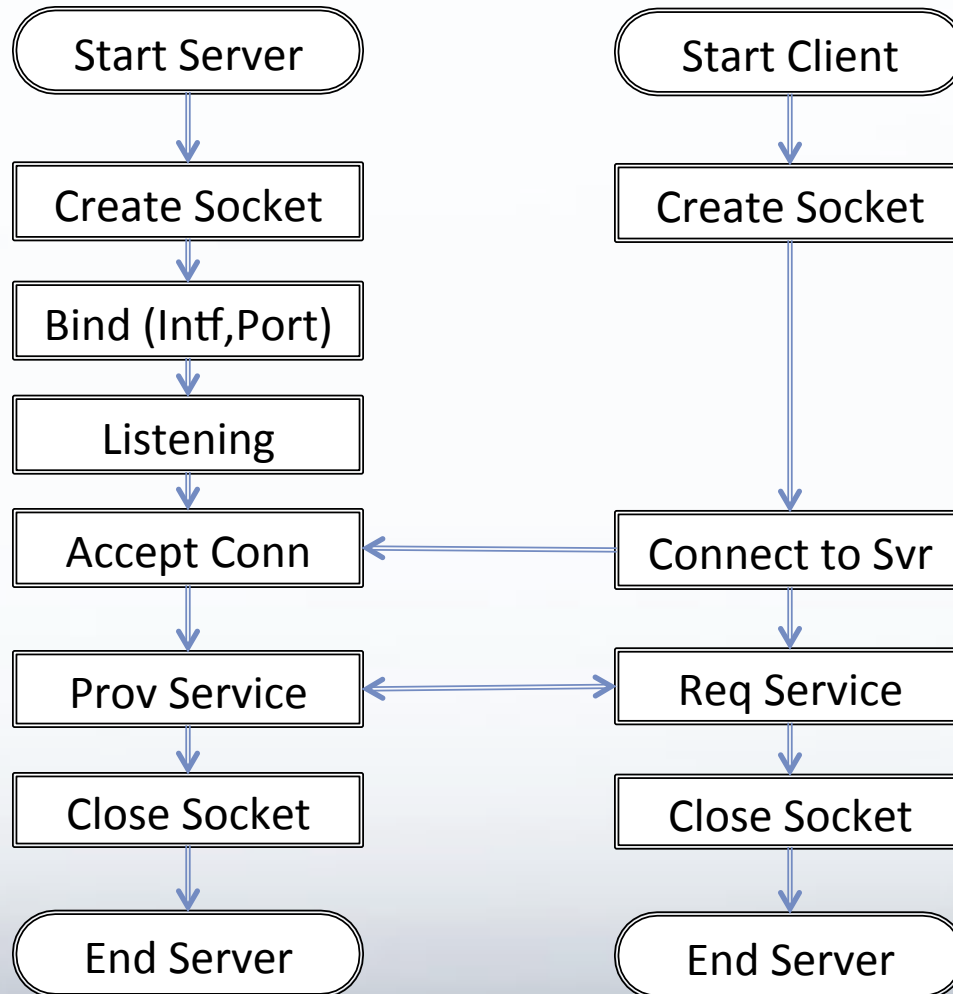
6. val = struct.pack("!h", 1)
7. s.sendto(val, (HOST, PORT))

8. data, (HOST, PORT) = s.recvfrom(1024)
9. print "Received %d" %
    (struct.unpack("!h", data))

10. s.close()
```

Connection-based communication

The flow chart



Connection demand

- In a connection-based communication, after binding the socket to an interface and port, the server must listen for incoming connection.
 - `socket.listen(backlog)`. *backlog* represents the number of incoming connection that can be queued at any time
- After listening for incoming connection, connection requests should be accepted
 - `socket.accept()`. `accept()` returns a pair *conn*, *address*, where
 - *conn* is a new socket object that the server will use to communicate with the remote host
 - *address* is the address of the remote host. The address format depends on the address family type
- The client connect to the server with the `connect()` method.
 - `socket.connect(address)`. *address* is the address of the remote host. The address format depends on the address family type

Sending data

- To send data, you can use
 - `socket.senda(string[, flags])`
 - Returns the number of bytes sent. The application must verify that whole data has been sent, or retry the data transmission if needed
 - `socket.sendall(string[, flags])`
 - Send all data unless an error occur
 - It returns “None” on success. Otherwise, an exception is raised and there is no way to know how many data has been sent
- For both methods, the optional *flags* argument can be used to execute special sending methods (e.g. out-of-band data)


Receiving data

- To receive data through a connected socket you can use `socket.recv(bufsize[, flags])`
 - `Bufsize` represents the maximum amount of bytes to read
 - `Flags` can be used to perform “special” readings
 - It returns a string which is actually the data read

Connection-based communication

Server

```
1. import socket
2. HOST = ''      # any available interf
3. PORT = 5000    # Arbitrary non-priv port
4. s = socket.socket(socket.AF_INET,
                    socket.SOCK_STREAM)
5. s.bind((HOST, PORT))
6. s.listen(1)
7. conn, addr = s.accept()
8. while 1:
9.     data = conn.recv(1024)
10.    if not data: break
11.    conn.sendall("Hi!")
12. conn.close()
13. s.close()
```



Client

```
1. import socket
2. HOST = '127.0.0.1'  # The
   remote host
3. PORT = 5000         # The
   remote port
4. s =
   socket.socket(socket.AF_INET,
                 socket.SOCK_STREAM)
5. s.connect((HOST, PORT))
6. val = "Hello!"
7. s.sendall(val)
8. data = s.recv(1024)
9. s.close()
10. print "Received: %s" %(data)
```

Server styles

- Iterating server
 - Only one socket is opened at a time
 - Clients are accepted one after the other
 - Slow service
- Forking server
 - With fork
 - After accept, the server creates a subprocess which will provide the service
 - The subprocess is a copy of the parent process
 - The used memory space is doubled
- With POSIX Threads
 - The same memory space is shared between all the threads
 - Special attention must be taken to avoid race conditions
- Concurrent single server
 - Uses select to simultaneously wait over the whole set of opened socket IDs
 - The main process is waken up when new data arrives
 - There is no context switching, but it cannot benefit from multiprocessors

Some thoughts about Multiprocessing

Multiprocessing in Linux / Unix-like systems

- Quick overview of sub-process creation and termination
- One can use the multiprocessing Python package to create and handle sub-processes
 - High-level API
 - Let's play with the OS services to understand the concepts
- `os.fork()` spawns a child process
 - Returns 0 in the child process
 - Returns the PID of the child in the parent process
 - In case of error, an `OSError` exception is raised

Example 1 – No synch between processes

```
1. import os
2. import time

3. def testdelay():
4.     for i in range(0,5):
5.         time.sleep(2)
6.         print "child is
   running... %d" %(i)
7.     os._exit(0)

8. pid = os.fork()
9. if pid == 0:
10.     testdelay()

11. print "parent is exiting..."
12. exit(0)
```

parent is exiting...
child is running... 0
child is running... 1
child is running... 2
child is running... 3
child is running... 4

Example 2- waiting for child termination

```
1. import os
2. import time

3. def testdelay():
4.     for i in range(0,5):
5.         time.sleep(2)
6.         print "child is running... %d" %
   (i)
7.     os._exit(0)

8. pid = os.fork()
9. if pid == 0:
10.     testdelay()
11. else:
12.     status = os.wait()
13.     print status

14. print "parent is exiting..."
15. exit(0)
```

child is running... 0
child is running... 1
child is running... 2
child is running... 3
child is running... 4
(59692, 0)
parent is exiting...

Example 3 – child exits faster than the parent process

```
1. import os
2. import time

3. def testdelay():
4.     for i in range(0,5):
5.         time.sleep(2)
6.         print "parent is
   running... %d" %(i)

7. pid = os.fork()
8. if pid == 0:
9.     print "child exits..."
10.    os._exit(0)
11. else:
12.     testdelay()
13.     status = os.wait()
14.     print status

15. print "parent is exiting..."
16. exit(0)
```

- Output
child exits...
parent is running... 0
parent is running... 1
parent is running... 2
parent is running... 3
parent is running... 4
(59692, 0)
parent is exiting...

- Process table status after child finishes but before parent finishes

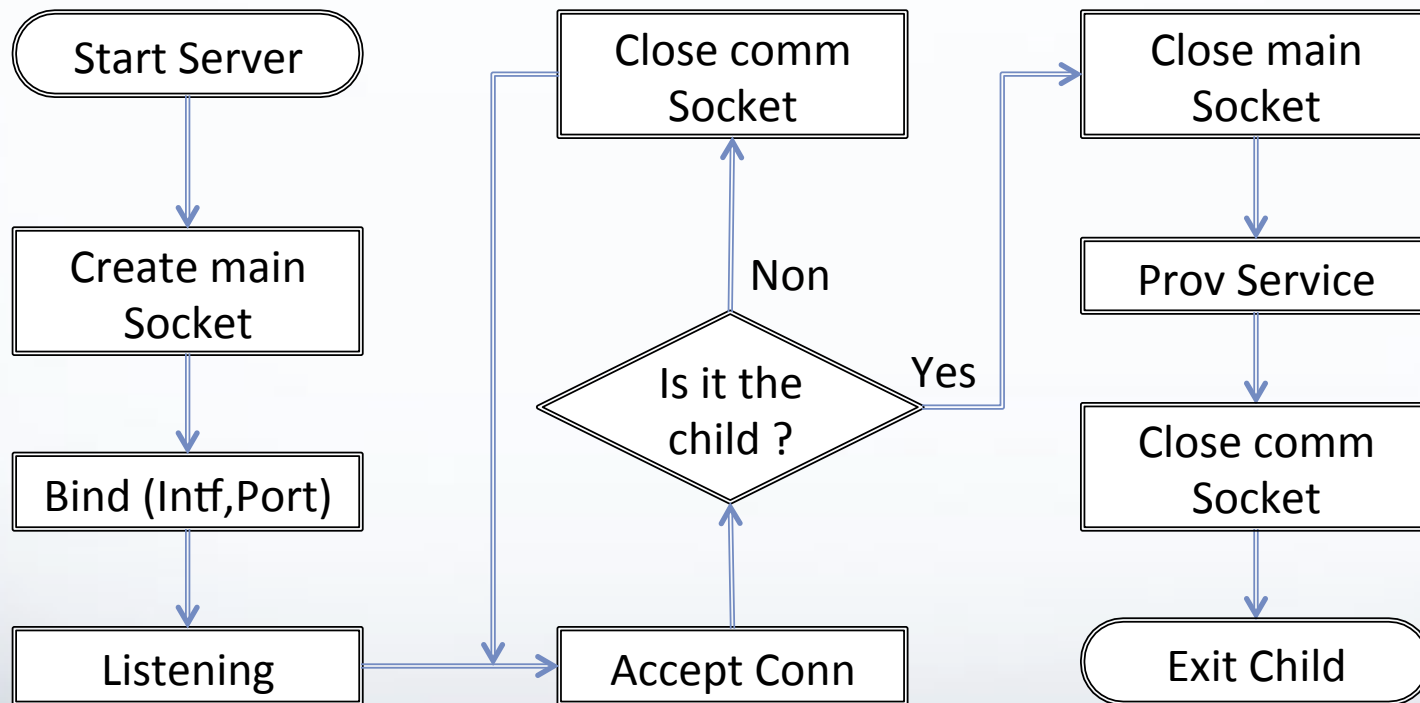
59691	ttys002	0:00.03	python test-fork.py
59692	ttys002	0:00.00	(Python)

↑
Zombie process

SIGCHLD and SIG_IGN

- Upon exit, a child process reports its exit code to its parent
- While the process parent doesn't read the exit status of the child process, this last is keep in the process table
 - Leading to a so-called zombie process
 - Refer to `wait()`, `waitpid()`
- In Linux, Unix-like systems, whenever something interesting happens to a forked off child, the parent process receives a SIGCHLD signal
- By default, SIGCHLD is ignored
- To avoid zombie process, the parent should handle the SIGCHLD signal. Ex.
 - `signal(signal.SIGCHLD, signal.SIG_IGN)`

Multi-process server



Multicast

Building a Multicast sender

- One-to-many communication
 - It is not broadcast
 - Multicast groups are identified by Class D IP addresses (224.0.0.1 → 239.255.255.255)
- Create your socket: SOCK_DGRAM or SOCK_STREAM ?
- Bind is optional. You should specify however on which interface multicast messages will be sent. Eg.
 - `s.setsockopt(socket.IPPROTO_IP, socket.IP_MULTICAST_IF, socket.INADDR_ANY)`
- Send the packet to the multicast address

Building a Multicast receiver

- Before binding your socket, it's used to reuse the local address, so multiple applications can use at the same time that multicast port
 - `s.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR, 1)`
- After binding, the receiver process must ask to the system to indeed receive any message addressed to the targeting multicast group and send it up to the application
 - `mreq = struct.pack("4sl", socket.inet_aton("226.1.1.1"), socket.INADDR_ANY)`
 - `s.setsockopt(socket.SOL_IP, socket.IP_ADD_MEMBERSHIP, mreq)`
- Before closing your socket, remove your membership (same operation as above, but with the **IP_DROP_MEMBERSHIP** option)