

Algorithms & Data Structures


Lesson 13: Minimum Spanning Trees

Marc Gaetano

Edition 2017-2018

Spanning Trees

- A simple problem: Given a *connected* undirected graph $\mathbf{G}=(\mathbf{V},\mathbf{E})$, find a minimal subset of edges such that \mathbf{G} is still connected
 - A graph $\mathbf{G2}=(\mathbf{V},\mathbf{E2})$ such that $\mathbf{G2}$ is connected and removing any edge from $\mathbf{E2}$ makes $\mathbf{G2}$ disconnected



Observations

1. Any solution to this problem is a tree
 - Recall a tree does not need a root; just means acyclic
 - For any cycle, could remove an edge and still be connected
2. Solution not unique unless original graph was already a tree
3. Problem ill-defined if original graph not connected
 - So $|E| \geq |V|-1$
4. A tree with $|V|$ nodes has $|V|-1$ edges
 - So every solution to the spanning tree problem has $|V|-1$ edges

Motivation

A **spanning tree** connects all the nodes with as few edges as possible

- Example: A “phone tree” so everybody gets the message and no unnecessary calls get made
 - Bad example since would prefer a balanced tree

In most compelling uses, we have a *weighted* undirected graph and we want a tree of least total cost

- Example: Electrical wiring for a house or clock wires on a chip
- Example: A road network if you cared about asphalt cost rather than travel time

This is the **minimum spanning tree** problem

- Will do that next, after intuition from the simpler case

Two Approaches

Different algorithmic approaches to the spanning-tree problem:

1. Do a graph traversal (e.g., depth-first search, but any traversal will do), keeping track of edges that form a tree
2. Iterate through edges; add to output any edge that does not create a cycle

Spanning tree via DFS


```
spanning_tree(Graph G) {  
    for each node i: i.marked = false  
    for some node i: f(i)  
}  
  
f(Node i) {  
    i.marked = true  
    for each j adjacent to i:  
        if(!j.marked)  
            add(i,j) to output  
            f(j) // DFS  
}
```

Correctness: DFS reaches each node. We add one edge to connect it to the already visited nodes. Order affects result, not correctness.

Time: $O(|E|)$

Example

Stack
f(1)




Output:

Example

Stack (bottom)

f(1)

f(2)



Output: (1,2)


Example

Stack (bottom)

$f(1)$

$f(2)$

$f(7)$



Output: (1,2), (2,7)

Example


Stack (bottom)

f(1)

f(2)

f(7)

f(5)



Output: (1,2), (2,7), (7,5)

Example

Stack (bottom)


f(1)

f(2)

f(7)

f(5)

f(4)



Output: (1,2), (2,7), (7,5), (5,4)

Example

Stack (bottom)

f(1)


f(2)

f(7)

f(5)

f(4)

f(3)



Output: (1,2), (2,7), (7,5), (5,4),(4,3)

Example

Stack (bottom)

f(1)


f(2)

f(7)

f(5)

f(4) f(6)

f(3)



Output: (1,2), (2,7), (7,5), (5,4), (4,3), (5,6)

Example

Stack (bottom)

f(1)


f(2)

f(7)

f(5)

f(4) f(6)

f(3)



Output: (1,2), (2,7), (7,5), (5,4), (4,3), (5,6)

Second Approach

Iterate through edges; output any edge that does not create a cycle

Correctness (hand-wavy):

- Goal is to build an acyclic connected graph
- When we add an edge, it adds a vertex to the tree
 - Else it would have created a cycle
- The graph is connected, so we reach all vertices


Efficiency:

- Depends on how quickly you can detect cycles
- Reconsider after the example

Example

Edges in some arbitrary order:

(1,2), (3,4), (5,6), (5,7), (1,5), (1,6), (2,7), (2,3), (4,5), (4,7)




Output:

Example

Edges in some arbitrary order:

(1,2), (3,4), (5,6), (5,7), (1,5), (1,6), (2,7), (2,3), (4,5), (4,7)




Output: (1,2)

Example

Edges in some arbitrary order:

(1,2), (3,4), (5,6), (5,7), (1,5), (1,6), (2,7), (2,3), (4,5), (4,7)




Output: (1,2), (3,4)

Example

Edges in some arbitrary order:

(1,2), (3,4), (5,6), (5,7), (1,5), (1,6), (2,7), (2,3), (4,5), (4,7)




Output: (1,2), (3,4), (5,6),

Example

Edges in some arbitrary order:

(1,2), (3,4), (5,6), (5,7), (1,5), (1,6), (2,7), (2,3), (4,5), (4,7)




Output: (1,2), (3,4), (5,6), (5,7)

Example

Edges in some arbitrary order:

(1,2), (3,4), (5,6), (5,7), (1,5), (1,6), (2,7), (2,3), (4,5), (4,7)




Output: (1,2), (3,4), (5,6), (5,7), (1,5)

Example

Edges in some arbitrary order:

(1,2), (3,4), (5,6), (5,7), (1,5), (1,6), (2,7), (2,3), (4,5), (4,7)




Output: (1,2), (3,4), (5,6), (5,7), (1,5)

Example

Edges in some arbitrary order:

(1,2), (3,4), (5,6), (5,7), (1,5), (1,6), (2,7), (2,3), (4,5), (4,7)




Output: (1,2), (3,4), (5,6), (5,7), (1,5)

Example

Edges in some arbitrary order:

(1,2), (3,4), (5,6), (5,7), (1,5), (1,6), (2,7), (2,3), (4,5), (4,7)



Can stop once we have $|V|-1$ edges

Output: (1,2), (3,4), (5,6), (5,7), (1,5), (2,3)

Cycle Detection

- To decide if an edge could form a cycle is $O(|V|)$ because we may need to traverse all edges already in the output
- So overall algorithm would be $O(|V||E|)$
- But there is a faster way: use union-find
 - Initially, each item is in its own 1-element set
 - Union sets when we add an edge that connects them
 - Stop when we have one set
 - Explain in next lesson

Summary So Far

The **spanning-tree problem**

- Add nodes to partial tree approach is $O(|E|)$
- Add acyclic edges approach is *almost* $O(|E|)$
 - Using union-find “as a black box”

But really want to solve the **minimum-spanning-tree problem**

- Given a weighted undirected graph, give a spanning tree of minimum weight
- Same two approaches will work with minor modifications
- Both will be $O(|E| \log |V|)$

Getting to the Point

Algorithm #1

Prim's Algorithm for Minimum Spanning Tree
is

Exactly our 1st approach to spanning tree
but process crossing edges in cost order

Algorithm #2

Kruskal's Algorithm for Minimum Spanning Tree
is

Exactly our 2nd approach to spanning tree
but process edges in cost order

Prim's Algorithm Idea


Idea:

- grow the MST starting with no edge
- mark vertices connected to the MST
- add an edge to the MST by picking the smallest weight edge among the *crossing* edges (*crossing edge*: edge with a vertex marked and a vertex not marked)


Algorithm pseudo-code

1. Set all vertices as *unmarked*
2. Set $S = \{ \}$, the set of edges of the MST
3. Set $C = \{ \}$, the set of crossing edge ((u, v) is a crossing edge iff \underline{u} is marked and v is unmarked)
4. Choose any node u
 - a) Mark \underline{u}
 - b) For each edge $e = (\underline{u}, v)$, () add e to C
5. While there are unmarked vertices in the graph
 - a) Select the crossing edge $e = (\underline{a}, b)$ with lowest cost
 - b) Add e to S
 - c) Mark \underline{b}
 - d) For each edge $e' = (\underline{b}, c)$ (c **not** marked), label e' “crossing”


Prim example




Prim example




Prim example




Prim example




Prim example




Prim example




Prim example



Prim example



Prim example



Prim's analysis

- Correctness
 - Invariant: S is a MST of the subgraph induced by the marked vertices
 - Variant: $|U|$ decrease down to 0 (U is the set of *unmarked* vertices)
- Run-time complexity
 - Sort the set of edges ($O(|E| \log |E|)$) and pick each edge in increasing order of weight ($O(|E|)$) = $O(|E| \log |E|)$
 - Somehow (non asymptotically) better: $O(|E| \log |E|)$ using a heap to store the crossing edges

Kruskal's algorithm idea

Idea:

- grow a forest out of edges that do not grow a cycle, just like for the spanning tree problem.
- But now consider the edges in order by increasing weight


So:

- Sort edges: $O(|E| \log |E|)$
- Iterate through edges: $O(|E|)$
- Use union-find for cycle detection: $O(|E|)$ (next lesson)

Kruskal's pseudocode

1. Set $S = \{ \}$, the set of edges of the MST
2. Sort edges by weight
3. Put each vertex in its own **set**
4. While the number of **sets** > 1
 - pick next smallest edge $e = (u, v)$
 - if u and v are in different sets S_1 and S_2
 - add e to S
 - merge S_1 and S_2

Kruskal's example



Edges in sorted order:

1: (A,D), (C,D), (B,E), (D,E)

2: (A,B), (C,F), (A,C)

3: (E,G)

5: (D,G), (B,D)


6: (D,F)

10: (F,G)

MST: { }

Sets: {A} {B} {C} {D} {E} {F} {G}

Kruskal's example



Edges in sorted order:

1: (A,D), (C,D), (B,E), (D,E)

2: (A,B), (C,F), (A,C)

3: (E,G)

5: (D,G), (B,D)


6: (D,F)

10: (F,G)

MST: { (A,D) }

Sets: {A, D} {B} {C} {E} {F} {G}

Kruskal's example



Edges in sorted order:

1: (A,D), (C,D), (B,E), (D,E)

2: (A,B), (C,F), (A,C)

3: (E,G)

5: (D,G), (B,D)


6: (D,F)

10: (F,G)

MST: { (A,D), (C,D) }

Sets: {A, D, C} {B} {E} {F} {G}

Kruskal's example



Edges in sorted order:

1: (A,D), (C,D), (B,E), (D,E)

2: (A,B), (C,F), (A,C)

3: (E,G)

5: (D,G), (B,D)


6: (D,F)

10: (F,G)

MST: { (A,D), (C,D), (B,E) }

Sets: {A, D, C} {B, E} {F} {G}

Kruskal's example



Edges in sorted order:

1: (A,D), (C,D), (B,E), (D,E)

2: (A,B), (C,F), (A,C)

3: (E,G)

5: (D,G), (B,D)


6: (D,F)

10: (F,G)

MST: { (A,D), (C,D), (B,E), (D,E) }

Sets: {A, D, C, B, E} {F} {G}

Kruskal's example



Edges in sorted order:

1: (A,D), (C,D), (B,E), (D,E)

2: (A,B), (C,F), (A,C)

3: (E,G)

5: (D,G), (B,D)


6: (D,F)

10: (F,G)

MST: { (A,D), (C,D), (B,E), (D,E) }

Sets: {A, D, C, B, E} {F} {G}

Kruskal's example



Edges in sorted order:

1: (A,D), (C,D), (B,E), (D,E)

2: (A,B), (C,F), (A,C)

3: (E,G)

5: (D,G), (B,D)


6: (D,F)

10: (F,G)

MST: { (A,D), (C,D), (B,E), (D,E), (C,F) }

Sets: {A, D, C, B, E, F} {G}

Kruskal's example



Edges in sorted order:

1: (A,D), (C,D), (B,E), (D,E)

2: (A,B), (C,F), (A,C)

3: (E,G)

5: (D,G), (B,D)


6: (D,F)

10: (F,G)

MST: { (A,D), (C,D), (B,E), (D,E), (C,F) }

Sets: {A, D, C, B, E, F} {G}

Kruskal's example



Edges in sorted order:

1: (A,D), (C,D), (B,E), (D,E)

2: (A,B), (C,F), (A,C)

3: (E,G)

5: (D,G), (B,D)

6: (D,F)

10: (F,G)

MST: { (A,D), (C,D), (B,E), (D,E), (C,F), (E,G) }

Sets: {A, D, C, B, E, F, G}

Kruskal's analysis

Correctness:

- invariant: S is a MST of the sub-graph induced by the set of sets of vertices
- variant: either the number of sets or the number of non chosen edges is decreasing

Runtime complexity:

- Floyd's algorithm to build min-heap with edges $O(|E|)$
- Iterate through edges using **deleteMin** to get next edge: $O(|E| \log |E|)$
- Use union-find to manage the set of sets of vertices: $O(|E|)$
- often stop long before considering all edges