


Lecture 7

Binary Heaps


Heaps

A *binary min-heap* (or just *binary heap* or just *heap*) is:

- **Structure property:** A *complete* binary tree
- **Heap property:** The priority of every (non-root) node is greater than (or equal to) the priority of its parent. AKA the children are always greater than the parents.



Array Representation of Binary Trees



From node i :

left child: $i * 2$

right child: $i * 2 + 1$

parent: $i / 2$

(wasting index 0 is
convenient for the
index arithmetic)

implicit (array) implementation:

	A	B	C	D	E	F	G	H	I	J	K	L	
0	1	2	3	4	5	6	7	8	9	10	11	12	13

Heap Operations Runtimes

insert and **deleteMin** both $O(\log N)$


at worst case, the number of swaps you have to do is the height of the tree. The height of a complete tree with N nodes is $\log N$.

Intuition:

1 Node

2 Nodes

4 Nodes



2^0 Nodes

2^1 Nodes

2^2 Nodes

Judging the array implementation

Plusses:

- Less "wasted" space
 - Just index 0 and unused space on right
 - In conventional tree representation, one edge per node (except for root), so $n-1$ wasted space (like linked lists)
 - Array would waste more space if tree were not complete
- Multiplying and dividing by 2 is very fast (shift operations in hardware)
- Last used position is just index **size**

Minuses:

- Same might-be-empty or might-get-full problems we saw with stacks and queues (resize by doubling as necessary)

Plusses outweigh minuses: "this is how people do it"

Build Heap

- Suppose you have n items to put in a new (empty) priority queue
 - Call this operation **buildHeap**
- n distinct **inserts** works (slowly)
 - Only choice if ADT doesn't provide **buildHeap** explicitly
 - $O(n \log n)$
- Why would an ADT provide this unnecessary operation?
 - Convenience
 - Efficiency: an $O(n)$ algorithm called Floyd's Method
 - Common tradeoff in ADT design: how many specialized operations

Floyd's Method

Intuition: if you have a lot of values to insert all at once, you can optimize by inserting them all and then doing a pass for swapping

1. Put the n values anywhere to make a complete structural tree
2. Treat it as a heap and fix the heap-order property
 - Bottom-up: leaves are already in heap order, work up toward the root one level at a time

```
void buildHeap() {  
    for(i = size/2; i>0; i--) {  
        val = arr[i];  
        hole = percolateDown(i, val);  
        arr[hole] = val;  
    }  
}
```

Example


- Build a heap with the values:
12, 5, 11, 3, 10, 2, 9, 4, 8, 1, 7, 6

- Stick them all in the tree to make a valid structure

- In tree form for readability.

Notice:

- Purple for node values to fix (heap-order problem)
- Notice no leaves are purple
- Check/fix each non-leaf bottom-up (6 steps here)




Algorithm Example

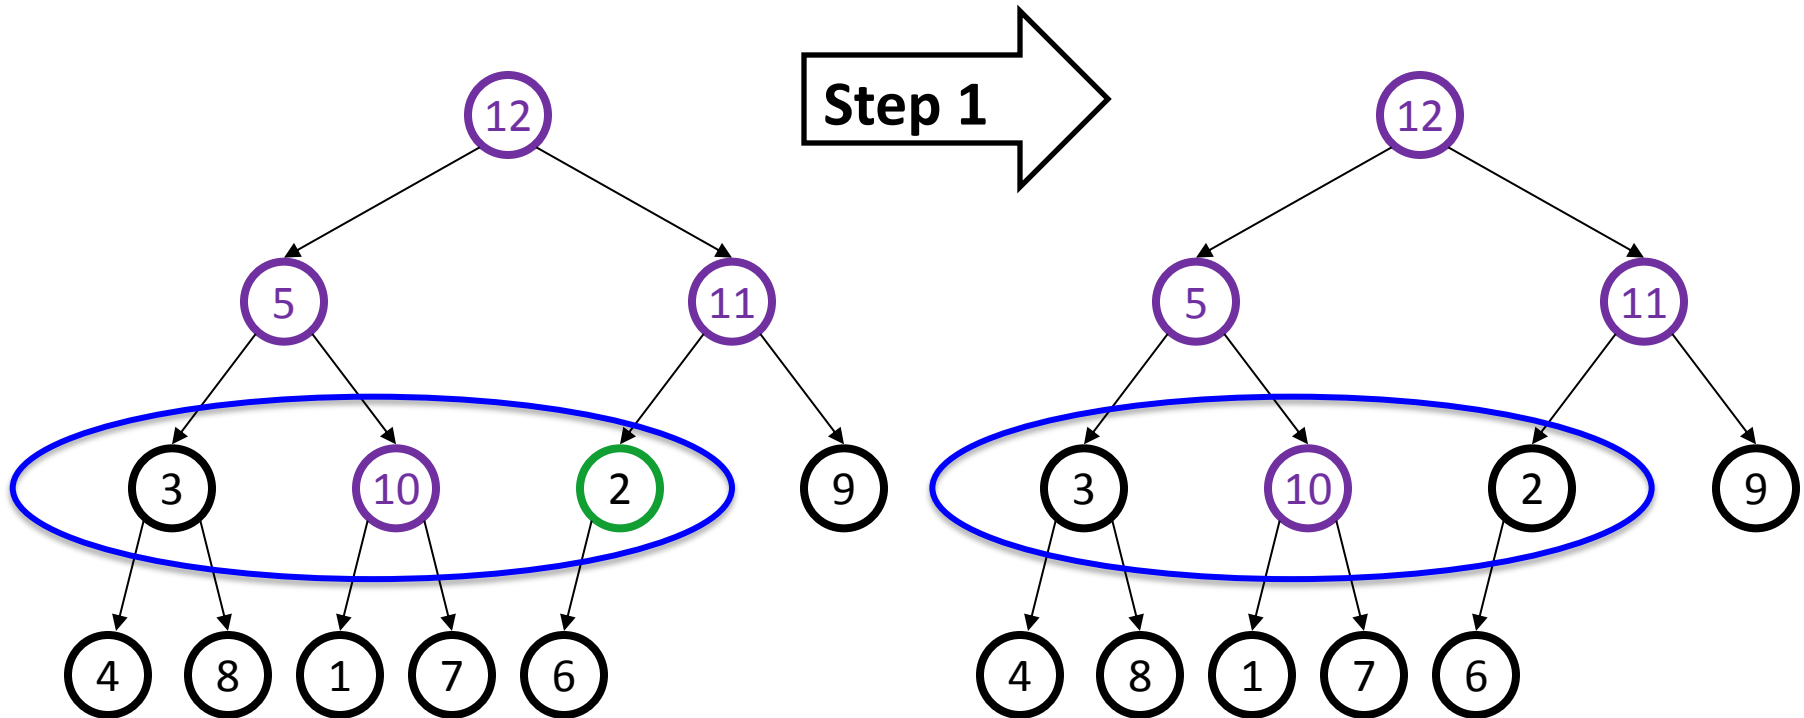
Purple shows the nodes that will need to be fixed.

We don't know which ones they are yet, so we'll traverse bottom up one level at a time and fix all the values.

Values to consider on each level circled in blue

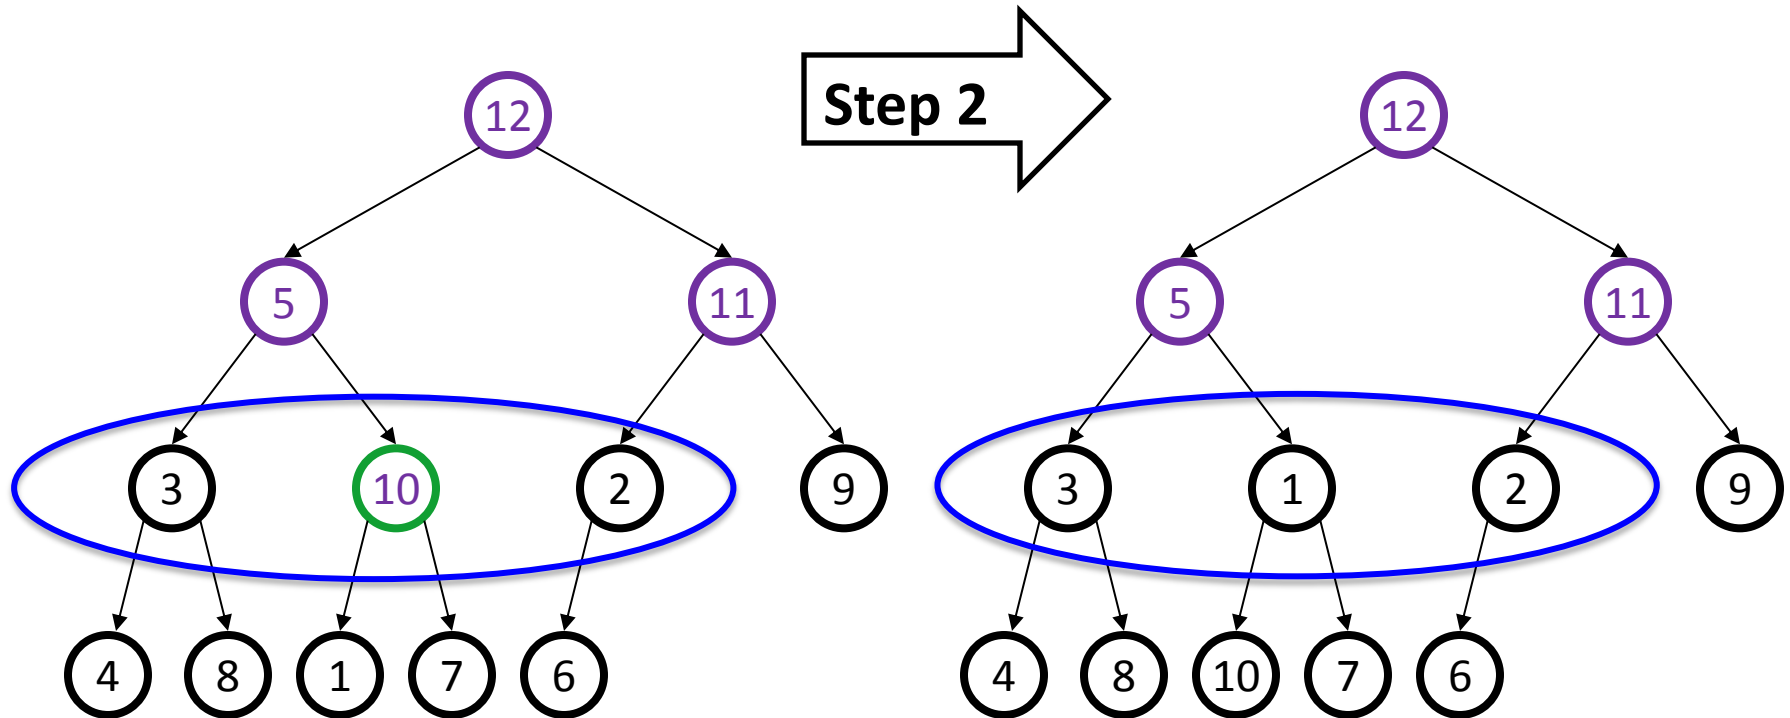


Algorithm Example



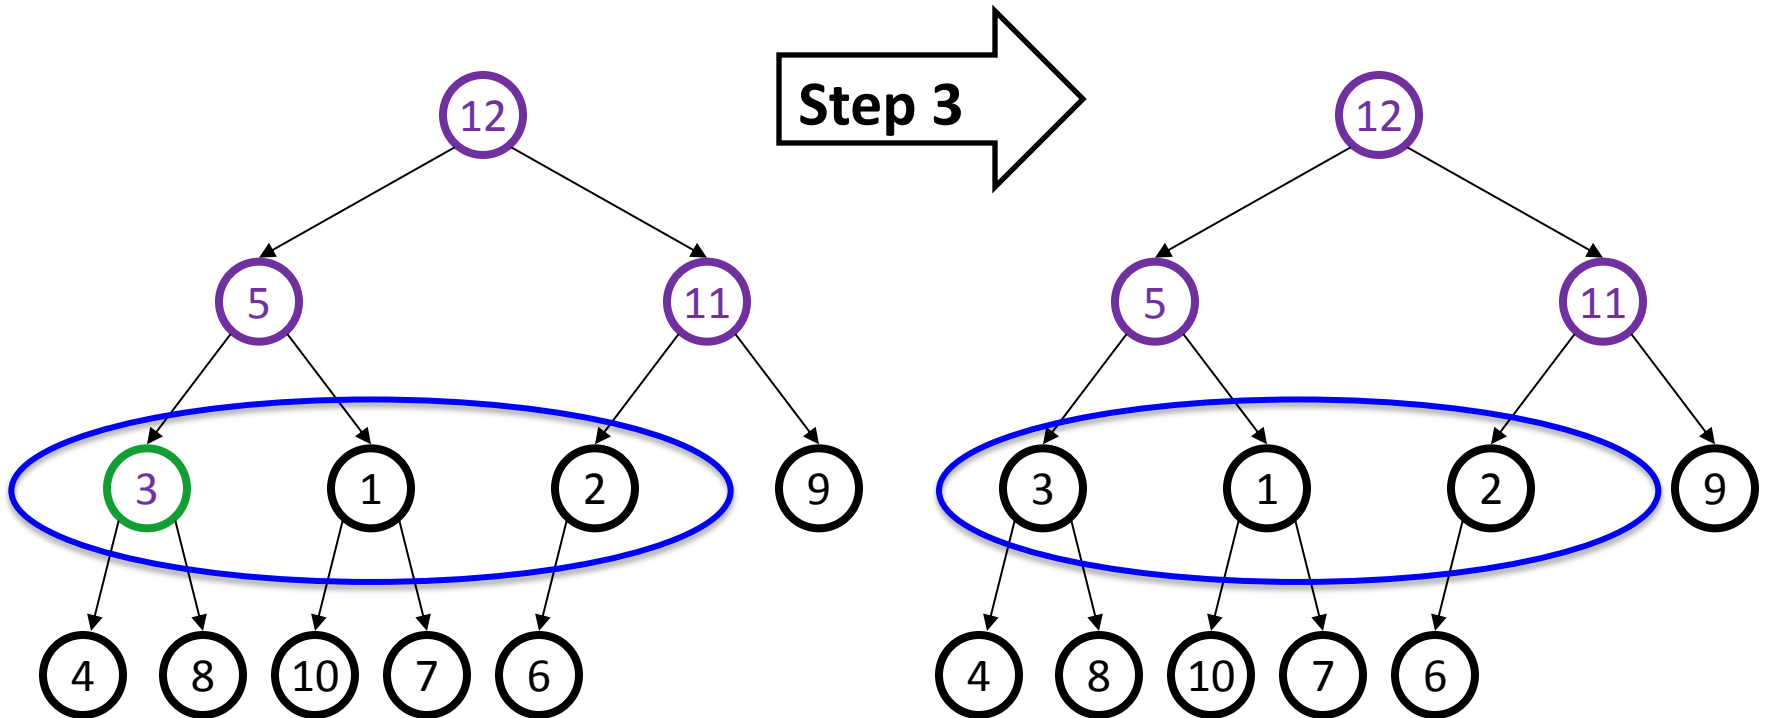
- Happens to already be less than it's child

Example



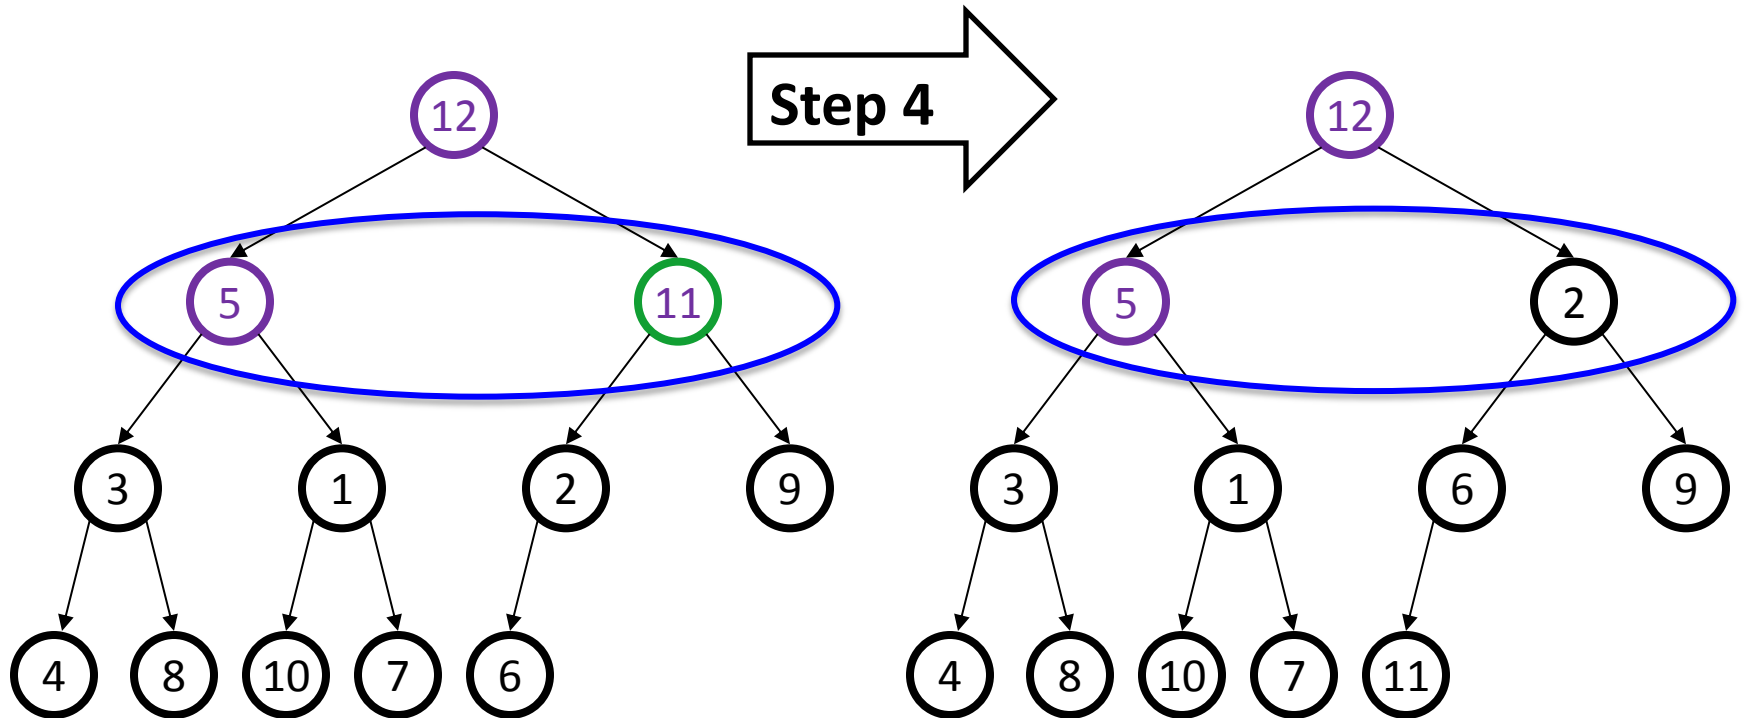
- Percolate down (notice that moves 1 up)

Example



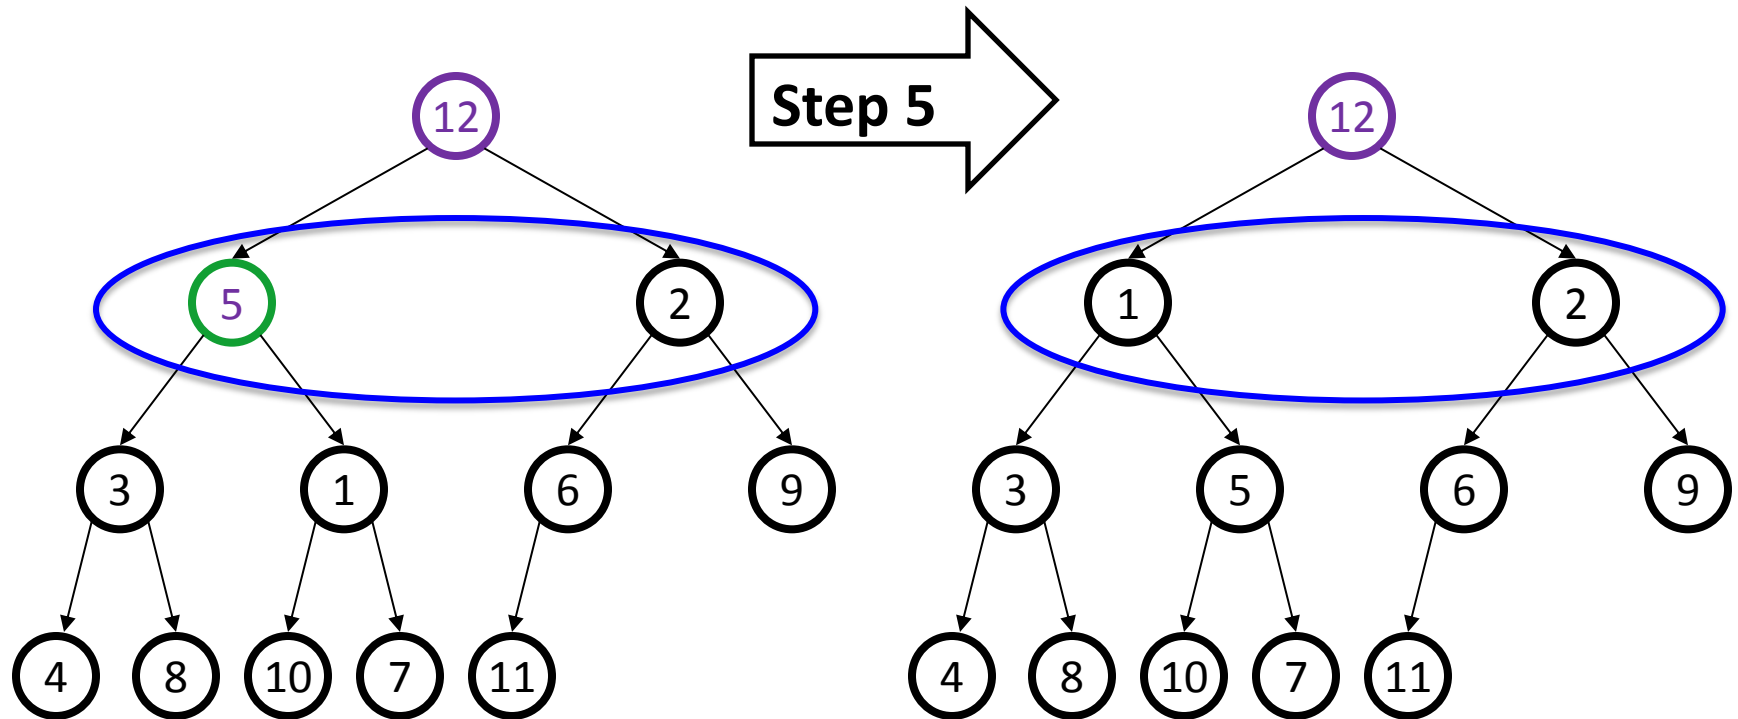
- Another nothing-to-do step

Example

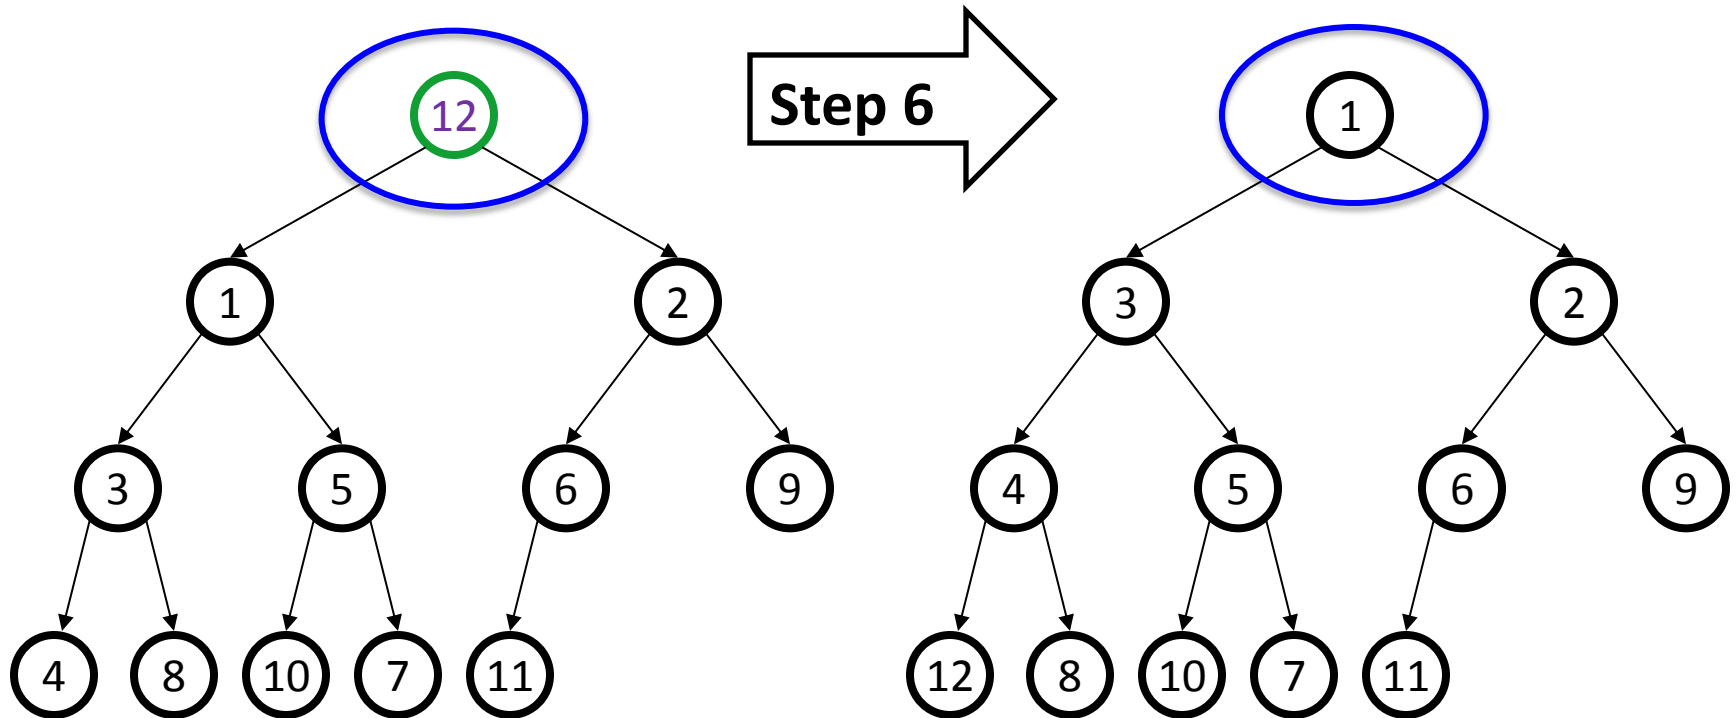


- Percolate down as necessary (steps 4a and 4b)

Example



Example



But is it right?

- “Seems to work”
 - Let’s *prove* it restores the heap property (correctness)
 - Then let’s *prove* its running time (efficiency)

```
void buildHeap() {  
    for(i = size/2; i>0; i--) {  
        val = arr[i];  
        hole = percolateDown(i, val);  
        arr[hole] = val;  
    }  
}
```


Correctness

```
void buildHeap() {  
    for(i = size/2; i>0; i--) {  
        val = arr[i];  
        hole = percolateDown(i, val);  
        arr[hole] = val;  
    }  
}
```

Loop Invariant: For all $j > i$, **arr[j]** is less than its children

- True initially: If $j > \text{size}/2$, then j is a leaf
 - Otherwise its left child would be at position $> \text{size}$
- True after one more iteration: loop body and **percolateDown** make **arr[i]** less than children without breaking the property for any descendants

So after the loop finishes, all nodes are less than their children

Efficiency

```
void buildHeap() {  
    for(i = size/2; i>0; i--) {  
        val = arr[i];  
        hole = percolateDown(i, val);  
        arr[hole] = val;  
    }  
}
```

Easy argument: **buildHeap** is $O(n \log n)$ where n is **size**

- **size/2** loop iterations
- Each iteration does one **percolateDown**, each is $O(\log n)$

This is correct, but there is a more precise (“tighter”) analysis of the algorithm...

Efficiency

```
void buildHeap() {  
    for(i = size/2; i>0; i--) {  
        val = arr[i];  
        hole = percolateDown(i, val);  
        arr[hole] = val;  
    }  
}
```

Better argument: **buildHeap** is $O(n)$ where n is **size**

- **size/2** total loop iterations: $O(n)$
- 1/2 the loop iterations percolate at most 1 step
- 1/4 the loop iterations percolate at most 2 steps
- 1/8 the loop iterations percolate at most 3 steps
- ...
- $((1/2) + (2/4) + (3/8) + (4/16) + (5/32) + \dots) < 2$ (page 4 of Weiss)
 - So at most $2(\text{size}/2)$ total percolate steps: $O(n)$

Lessons from **buildHeap**

- Without **buildHeap**, our ADT already let clients implement their own in $O(n \log n)$ worst case
 - Worst case is inserting better priority values later
- By providing a specialized operation internal to the data structure (with access to the internal data), we can do $O(n)$ worst case
 - Intuition: Most data is near a leaf, so better to percolate down
- Can analyze this algorithm for:
 - Correctness:
 - Non-trivial inductive proof using loop invariant
 - Efficiency:
 - First analysis easily proved it was $O(n \log n)$
 - Tighter analysis shows same algorithm is $O(n)$

Other function

- **merge**: given two priority queues, make one priority queue
 - How might you merge binary heaps:
 - If one heap is much smaller than the other?
 - If both are about the same size?
 - Different pointer-based data structures for priority queues support logarithmic time **merge** operation (impossible with binary heaps)
 - Leftist heaps, skew heaps, binomial queues
 - Worse constant factors
 - Trade-offs!