

Les données numériques: corrélation pour les images

Diane Lingrand



2022 - 2023

1 Correlation

2 Lab

1 Correlation

2 Lab

- For f and g discrete functions :




$$f \star g(x, y) = \sum_{u=-\infty}^{\infty} \sum_{v=-\infty}^{\infty} f(x, y)g(x - u, y - v)$$

- Applied to image I_1 by kernel k of shape $(2p + 1) \times (2q + 1)$:




$$\begin{aligned} I_2[x, y] &= I_1 \star k[x, y] \\ &= \sum_{i=0}^{2p} \sum_{j=0}^{2q} I_1[x + i - p, y + j - q]k[i][j] \end{aligned}$$

- same kernel applied to different locations on the image
- equivalent to convolution if the kernel is symmetric (horizontally and vertically).

Toy example



Toy example






- Cross-Correlation of image I_1 by kernel k of shape $(2p + 1) \times (2q + 1)$ ($N = (2p + 1)(2q + 1)$) :

$$I_2[x, y] = I_1 \star k[x, y]$$
$$= \frac{1}{\sigma_k \sigma_{I_{xy}} N} \sum_{i=0}^{2p} \sum_{j=0}^{2q} (I_1[x + i - p, y + j - q] - \mu_{I_{xy}})(K[i][j] - \mu_k)$$

- where :
 - μ_k is the mean of kernel values
 - σ_k is the standard deviation of kernel values
 - $\mu_{I_{xy}}$ (resp. $\sigma_{I_{xy}}$) is the mean (resp. standard deviation) of pixel values in the neighborhood of coordinates $(x; y)$. The dimensions of the neighborhood are those of the kernel.
- values between -1 and +1.

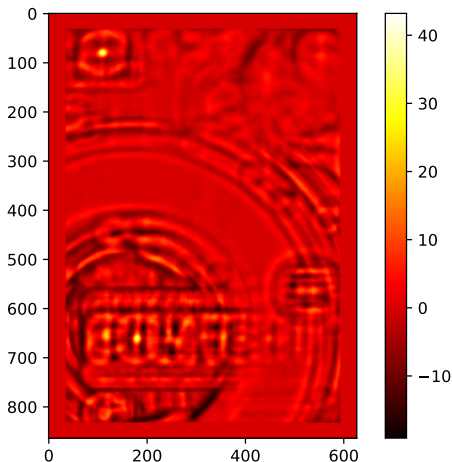
Toy example




Find a logo in an image



Find  in



Grey levels. Correlation results plotted using hot colormap. Max 0.999 at (82;110).



Superposition of grey image with highest values from the correlation map (> 0.5)

Naive implementation

```
def crossCorrelation(im, mo):
    hmo, wmo = mo.shape[:2]
    hmo2, wmo2 = hmo//2, wmo//2 # assuming mo of odd width and height
    him, wim = im.shape[:2]
    res = np.zeros(im.shape)
    mo = mo-np.mean(mo)
    stdMo = math.sqrt(np.mean(mo*mo))
    if stdMo == 0:
        return res
    mo /= stdMo
    for j in range(wmo2, wim-wmo2):
        for i in range(hmo2, him-hmo2):
            #mean of roi
            meanRoi = 0
            for l in range(wmo):
                for k in range(hmo):
                    meanRoi += im[i+k-hmo2][j+l-wmo2]
            meanRoi /= hmo*wmo
            # standard dev of roi
            stdRoi = 0
            for l in range(wmo):
                for k in range(hmo):
                    stdRoi += (im[i+k-hmo2][j+l-wmo2]-meanRoi)**2
            stdRoi = math.sqrt(stdRoi/hmo/wmo)
            if stdRoi == 0:
                res[i][j] = 0
            #cross correlation
            else:
                for l in range(wmo):
                    for k in range(hmo):
                        res[i][j] += (im[i+k-hmo2][j+l-wmo2]-meanRoi)*mo[k][l]
                res[i][j] /= stdRoi
    res /= hmo*wmo
    return res
```

Improve the implementation

- numpy functions are faster
 - compute mean and standard deviation using numpy
 - compute correlation as a pointwise multiplication in numpy
- test the previous example
 - and observe the difference

```
1 print(a, "\t\tmatrix a\n")
2 print(b, "\t matrix b\n")
3 print(a*b, "\t pointwise multiplication\n")
4 print(a@b, "\t matrix multiplication")
```

```
[[1 2 3]
 [4 5 6]
 [7 8 9]]
```

matrix a

```
[[ 10  20  30]
 [ 200 100 200]
 [1000 2000 3000]]
```

matrix b

```
[[ 10  40  90]
 [ 800 500 1200]
 [ 7000 16000 27000]]
```

pointwise multiplication

```
[[ 3410  6220  9430]
 [ 7040 12580 19120]
 [10670 18940 28810]]
```

matrix multiplication


Find **a**'s in

where :


- μ_k is the mean of kernel values
- σ_k is the standard deviation of kernel values
- $\mu_{l_{xy}}$ (resp. $\sigma_{l_{xy}}$) is the mean (resp. standard deviation) the neighborhood of coordinates $(x; y)$. The dimension neighborhood are those of the kernel.

Find letters in a text


correlation map



correlation values > 0.5
highlighted with yellow circles



localisation of max. correlations
on the original image



- Implement the correlation without and with normalisation
 - do not forget to use `numpy` for improvements
- Choose an image with different instances of logos. It can be a text with letters.
 - compute and visualise the correlation map
 - find the highest values (`numpy.argmax(condition)` can help)
 - redraw the highest values locations on the original image
 - for example, this was used on previous slide :

```
for (a,b) in lesA:  
    c = Circle((b,a), color='yellow',radius=14, linewidth=1, fill=False)  
    ax.add_patch(c)
```

1 Correlation

2 Lab

MNIST dataset

Dataset of 60000+10000 grayscale squared images (28x28).




Reduced version of MNIST

- 10 classes of gray images (28x28)
- only 2000 images in the train set (200 per class)
- only 1000 images in the test set (100 per class)
- How to get the files ?
 - download :
 - <https://www.i3s.unice.fr/~lingrand/redMNIST-x-train.bin>
 - <https://www.i3s.unice.fr/~lingrand/redMNIST-x-test.bin>
 - <https://www.i3s.unice.fr/~lingrand/redMNIST-y-train.bin>
 - <https://www.i3s.unice.fr/~lingrand/redMNIST-y-test.bin>
 - load in you python code :


```
import pickle
with open('redMNIST-x-train.bin', 'rb') as input:
    xTrain = pickle.load(input)
with open('redMNIST-x-test.bin', 'rb') as input:
    xTest = pickle.load(input)
with open('redMNIST-y-test.bin', 'rb') as input:
    yTest = pickle.load(input)
with open('redMNIST-y-train.bin', 'rb') as input:
    yTrain = pickle.load(input)
```

Looking at images

```
1 import random
2 n = random.randrange(0, len(yTrainr))
3 plt.imshow(xTrain[n], cmap=plt.cm.gray)
```



- Confusion matrix computed on the test set :



- accuracy = 95.1%

How to do that ?

- main idea :
 - the results of correlation/convolution of some filters will help to represent an image in a more compact and mean full than the $28 * 28 = 784$ pixel values
- which filters ?
 - can be handcrafted : $\cap, \cup, \bigcirc, \subset, \supset, |, -, /$
 - I can give you some other filters
- It is not necessary to compute corr./conv. at every position :
 - skip some rows and cols
 - for example : 1 over 3 cols and 1 over 3 rows
 - it will speed the computations
- normalisation
 - divide every MNIST digit by 255
 - use almost normalised filters
 - and thus skip the normalisation of correlation in order to speed
 - eventually, remove all negative values
- when each image is represented using some corr./conv. results
 - perform a classification :
 - using kNN
 - using another algorithm that we will see next week : logistic regression

Why this approach ?

- It is a simplified version of a well known deep neural network for image classification :
 - CNN = Convolutional Neural Network
 - except that a CNN :
 - will learn the coefficient of corr./conv.
 - may stack different series of corr./conv. filters
 - may use a non linear classification
 - may use different other tricks
- You will learn true CNN next year
- Here are some filters you could test :

