

Real-Time Operating Systems (RTOS) for sensors and actuators

B. Miramond

Polytech Nice Sophia

Some existing solutions



IoT OS and RTOS

- **Open-source**

TinyOS, RIOT, Contiki,
Mantis OS, Nano RK,
LiteOS, FreeRTOS, Apache
Mynewt, Zephyr OS,
Ubuntu Core 16 (Snappy),
ARM mbed, Yocto,
Raspbian

- **Commercial**

Android Things, Windows
10 IoT, WindRiver
VxWorks, **Micrium μ C/OS**,
Micro Digital SMX RTOS,
MicroEJ OS, Express Logic
ThreadX, TI RTOS,
Freescale MQX, Mentor
Graphics Nucleus RTOS,
Green Hills Integrity,
Particle

IoT OPERATING SYSTEMS

Which operating system(s) do you use for your IoT devices?



How to select one ?

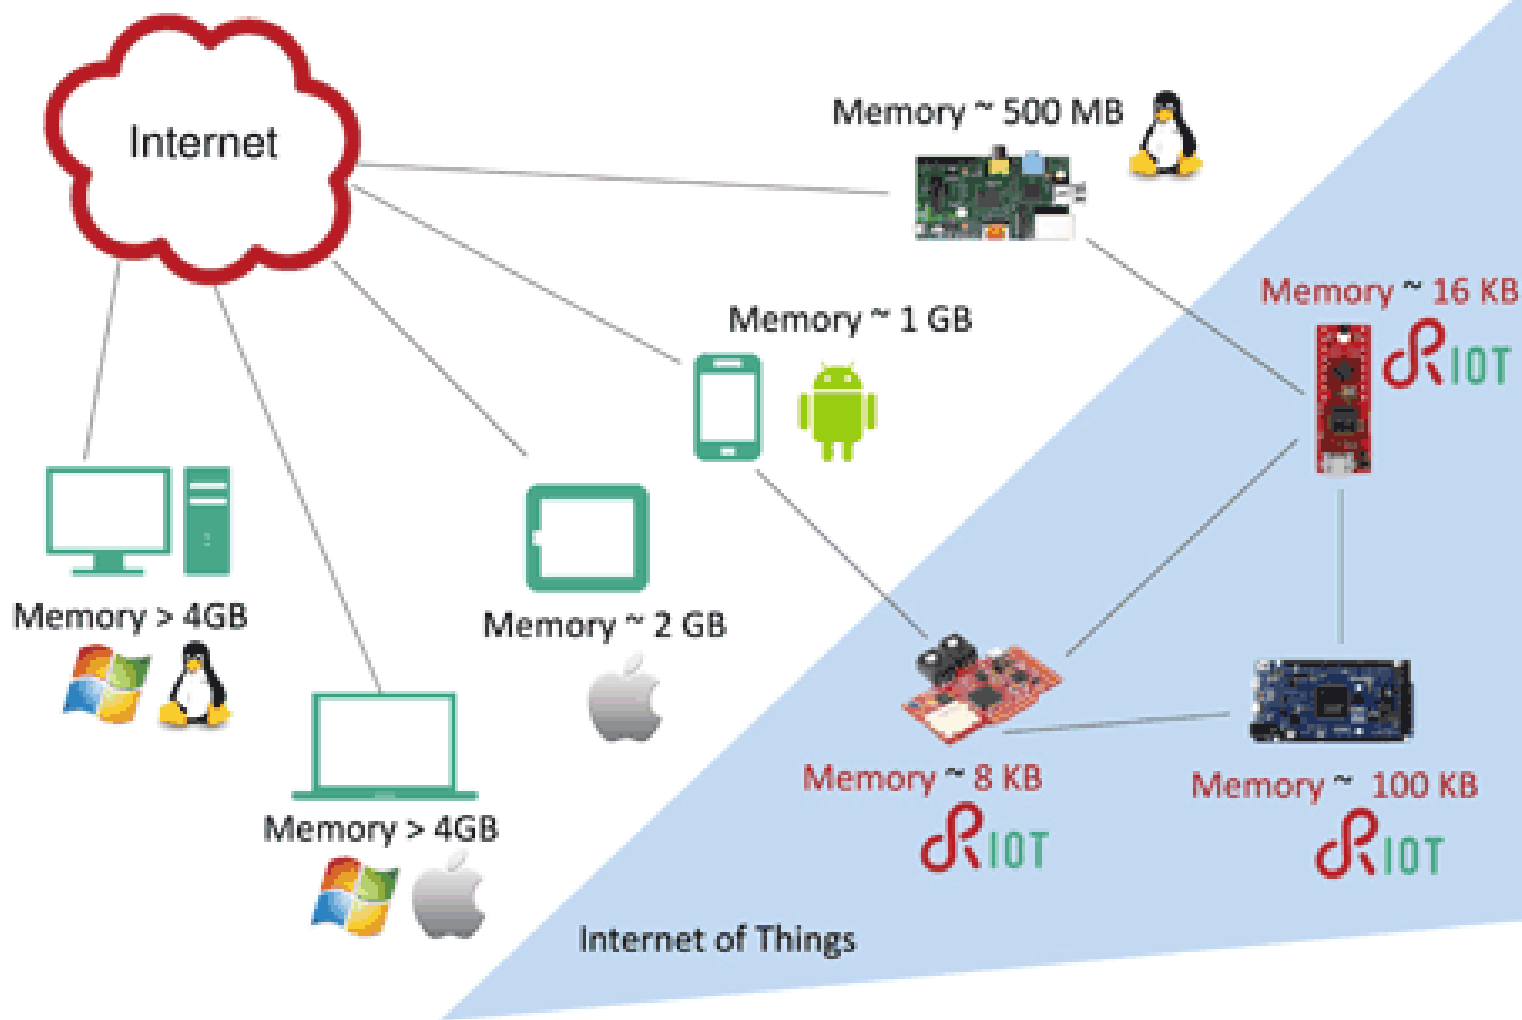
- Sensors are constrained systems that need:
 - Reduced memory footprint (ROM & RAM)
 - Low time-overhead (2 to 4 % of CPU time)
- But also
 - Reduced cost (royalties per unit)
 - Maintenance and support cost

How to select one ?

- **Footprint:** Since devices are constraint, we expect OS to have low memory, power and processing requirements. The overhead due to the OS should be minimal.
- **Portability:** OS isolates applications from the specifics of the hardware. Usually, OS is ported to different hardware platforms and interfaces to the board support package (BSP) in a standard way, such as using POSIX calls.
- **Modularity:** OS has a kernel core that's mandatory. All other functionality can be included as add-ons if so required by the application.
- **Connectivity:** OS supports different connectivity protocols, such as Ethernet, Wi-Fi, BLE, IEEE 802.15.4, and more.
- **Scalability:** OS must be scalable for any type of device. This means developers and integrators need to be familiar with only one OS for both nodes and gateways.
- **Reliability:** This is essential for mission-critical systems. Often devices are at remote locations and have to work for years without failure. Reliability also implies OS should fulfil certifications for certain applications.
- **Security:** OS has add-ons that bring security to the device by way of secure boot, SSL support, components and drivers for encryption.

Typical memory requirements

Source: Baccelli et al. 2015



Our use-case: real-time kernel uC/OS-II

- uC/OS-II (Micro Controller Operating System)
 - 1992 (J. Labrosse)
 - 50 (8-64 bits)
 - Portable
 - ROMable (6K-24K bytes) + 1K bytes RAM
 - Scalable
 - Preemptive fixed priorities
 - Deterministic
 - With or without MMU

ALTERA

ANALOG
DEVICES

ARM

Atmel

CYPRESS
PERFORM

FUJITSU

Infineon

Microsemi

NXP

RENESAS

SAMSUNG

SILICON LABS

life.augmented

TEXAS INSTRUMENTS

TOSHIBA
Leading Innovation >>>

XILINX

uc/OS-II vs. uc/OS-III

Feature	μC/OS-II	μC/OS-III
Release Date	1999–present	2009–present
Preemptive Multitasking	✓	✓
Maximum number of tasks	255	Unlimited
Number of tasks at each priority level	1	Unlimited
Round robin scheduling		✓

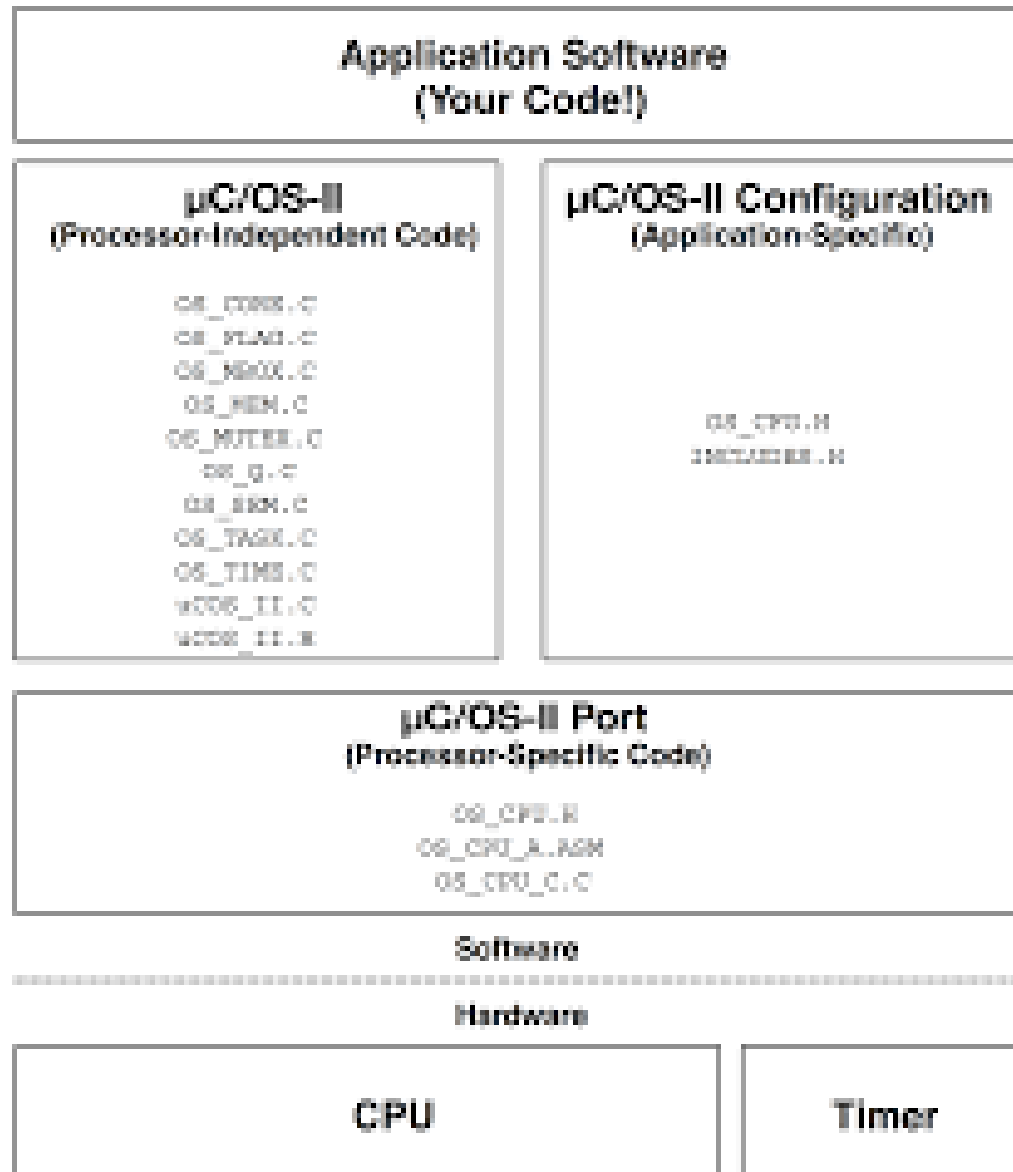
uC/OS-II Tasks

- Maximum 64 tasks
- Exclusive priorities
- No round robin (2 tasks don't have the same priority)
- Tasks can communicate through ISR (Interrupt Service Routine)

Services of the real-time kernel

- The uC/OS kernel provides the following services:
 - Scheduling
 - Change of context
 - Management of synchronization modes (semaphores, mutex)
 - Management of the means of communication (mailbox, queue)
 - Delays and timeout
- This gives a memory footprint of a few KB.

Structure of uC/OS-II




Memory footprint

- The total size of the application with an RTOS depends on:
 - Size of the application code
 - Data (RAM) and instruction (ROM) space used by the kernel
 - $\text{SUM}(\text{Task Stack} + \text{TCB}) // \text{Task Control Block}$
 - $\text{SUM}(\text{ECB}) / \text{Event Control Block}$
 - ISR routines
- For systems with little memory, you have to pay attention to the evolution of the stack size :
 - Table and local structures with functions
 - Nested function calls
 - Number of function arguments

Tasks under ucOS-II

- The uC/OS-II Scheduler manages the following tasks:
 - User tasks
 - Plus 2 system tasks :
 - Idle Task (OS_LOWEST_PRIO)
 - Statistics Task (OS_LOWEST_PRIO - 1)
- Since the priorities are coded as an INT8U, and the priorities are not shared (no round robin), the number of possible tasks is
$$64 - 2 = 62 \text{ user tasks.}$$

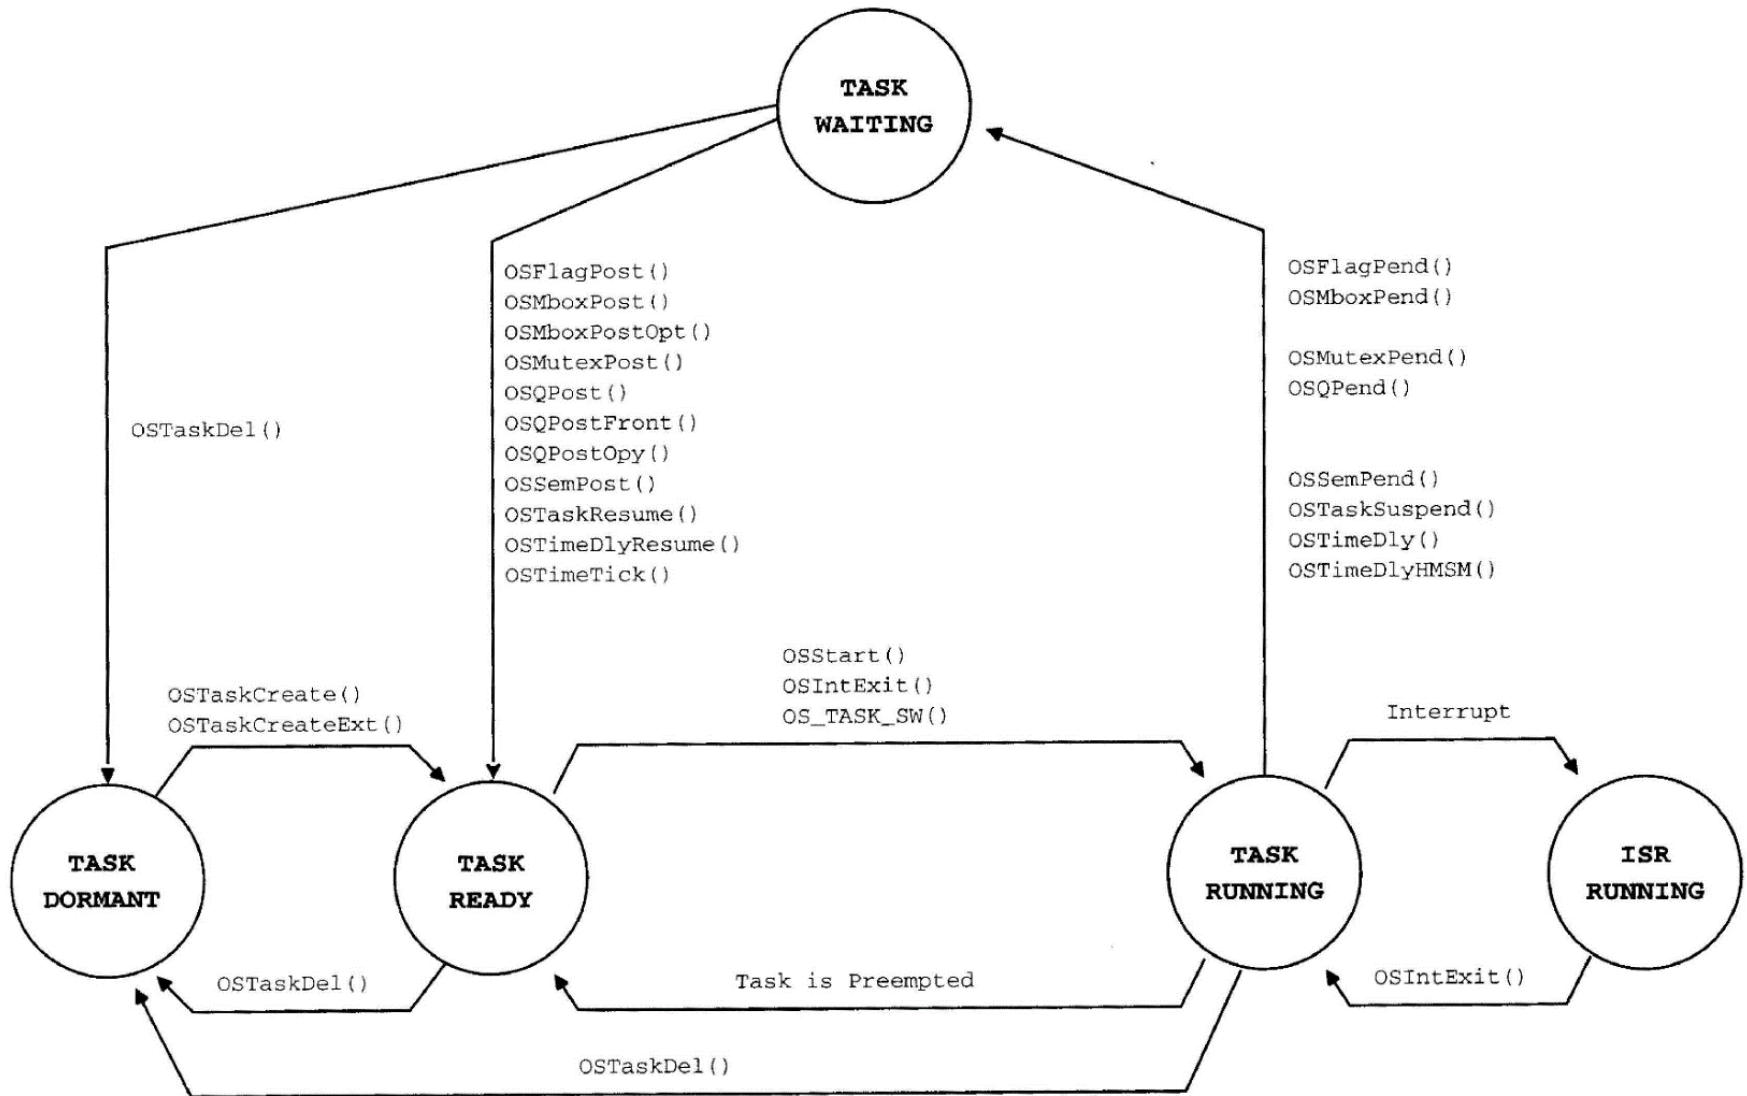
General state diagram



2 more states with uCOS-II

- **DORMANT :**
 - Before being created
 - After task destruction (OSTaskDel())
- **ISR :**
 - When an interrupt is triggered
 - On return, the current task can be pre-empted by a higher priority task (RDY state).
- **WAITING :**
 - Waiting for an event or timeout(OSTimeDly() or OS_X_Pend())
- **READY :**
 - Creation of tasks before multitasking is launched (OSStart())
 - Preempted
- **RUNNING :**
 - The only task using the processing resource (on MCU without multi-threading)

State diagram under ucOS



BUT WHAT IS A TASK ?

A task is managed through a data structure: TCB

The OS keeps the context of a task in a specific data structure called the **TCB**.

There is a change of context (save/restore) at each pre-emption or change of task:

- Processor registers
- General registers: PC, SP, section registers
- Runtime stack
- Task status
- ...

Task ID

State

CPU Context

Memory Context

Ressource

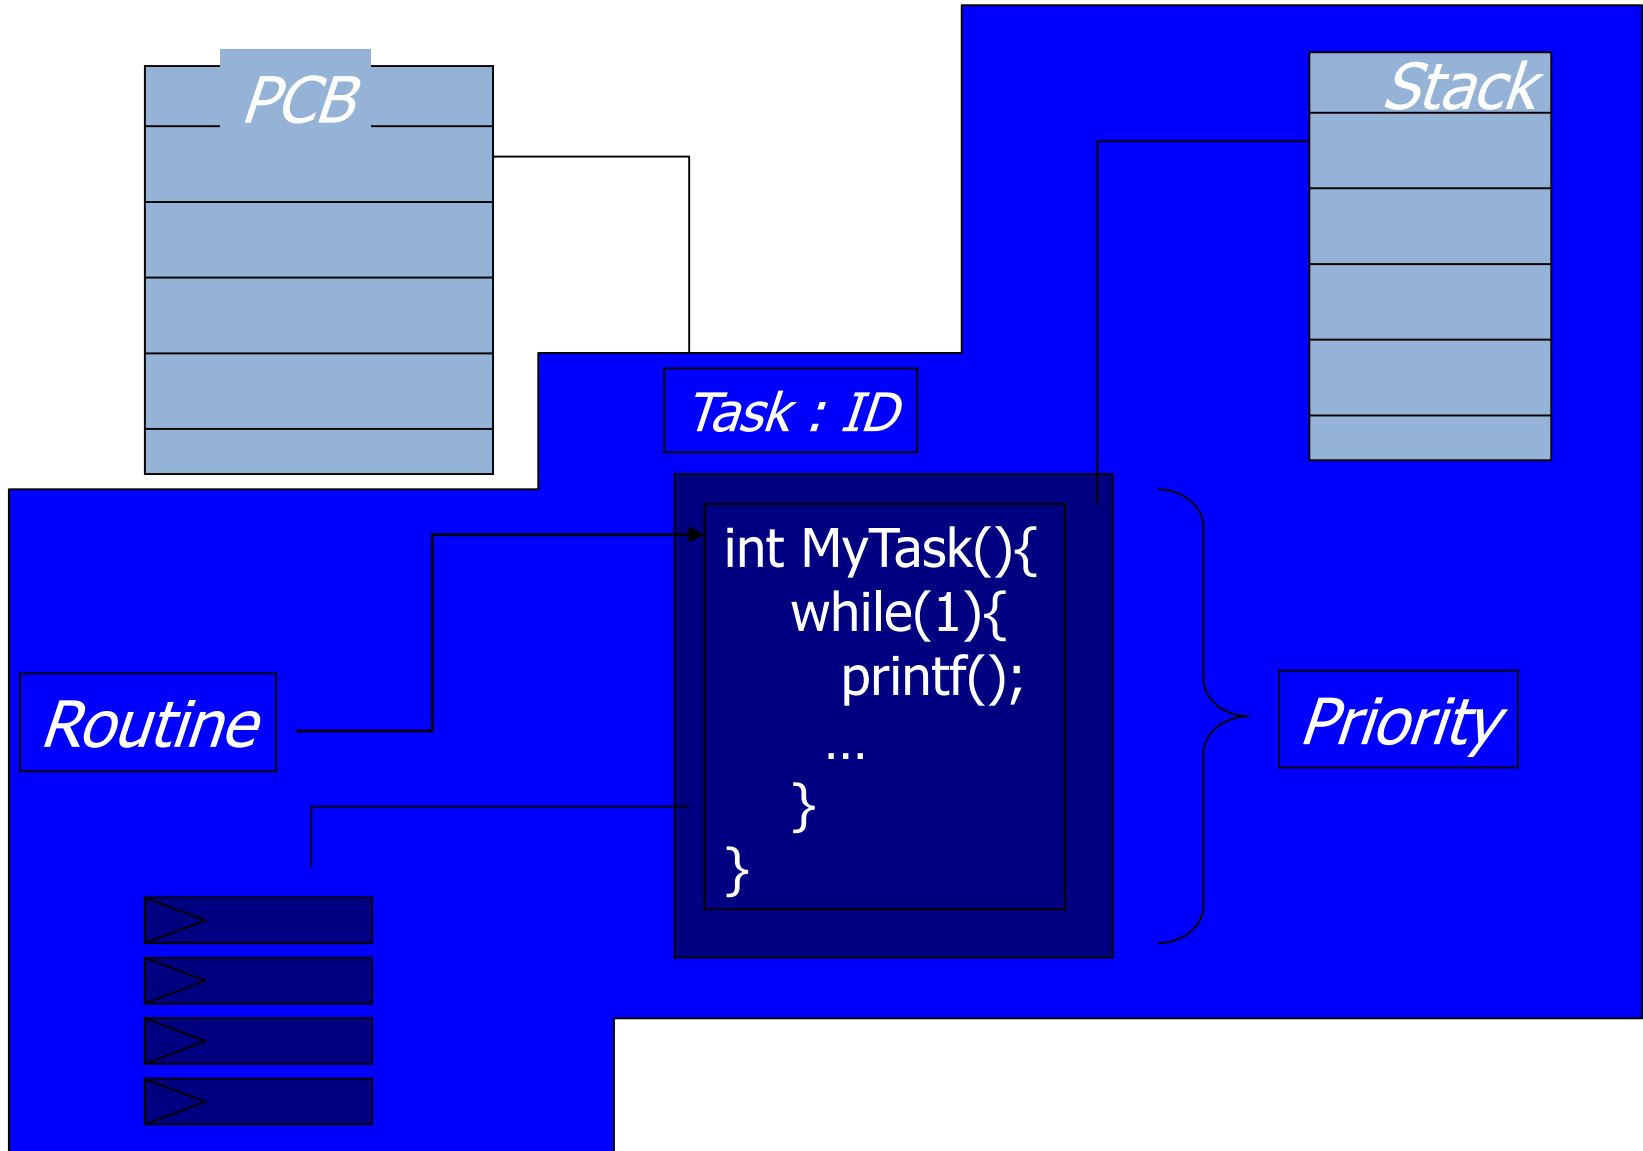
Scheduling data

User Information

More precisely

- The identifier of the process as it was assigned at its creation (an integer)
- One of the states of the process
- The value of the processor registers (PC, RI, SR...)
- The start and end addresses of the execution stack
- The open files,
- The synchronization services used
- The priority of the process, its queue...
- The CPU time used, the size of its stack, ...

Illustration of the task



uC-OS-II services API

- [Task Management](#)
- [Time Management](#)
- [Timer Management](#)
- [Event Control Blocks](#)
- [Semaphore Management](#)
- [Mutual Exclusion Semaphores](#)
- [Event Flag Management](#)
- [Message Mailbox Management](#)
- [Message Queue Management](#)
- [Memory Management](#)
- [Porting \$\mu\$ C/OS-II](#)

API documentation :

- <https://micrium.atlassian.net/wiki/spaces/osiidoc/overview>

Fonctions to create a tasks:

OSTaskCreate() and OSTaskCreateExt()

- Rule 1: Tasks must be created before launching the OSStart() function that launches the multi-tasking management mechanism.
- Rule 2: Ext version is an extended version used only for code analysis.
- The first one has 4 arguments, the second one 9.

Arguments of the task creation function

- void (*task)(void*pd): pointer on the task function
- Void * pdata : payload pointer used at the creation
- OS_STK *ptos : pointer on top of stack
- INT8U prio : priority of the task (limitations to 64 values)
- INT16U id : task ID (extension of the previous limitation to 64 values, else id = prio)
- OS_STK *pbos : Bottom-of-stack
- INT32U stk_size : stack size (for stack checking).
- void *pext : user data extension
- INT16U opt : uCOS_ii.h contains the list of options (OS_TASK_OPT_STK_CHK, OS_TASK_OPT_TSK_CLR, OS_TASK_OPT_SAVE_FP...). Each constant is a binary flag.

TCB

OS_STK = by default is 2 bytes long

```

*****
*/

typedef struct os_tcb {
    OS_STK          *OSTCBStkPtr;          /* Pointer to current top of stack          */
    #if OS_TASK_CREATE_EXT_EN > 0
        void          *OSTCBExtPtr;        /* Pointer to user definable data for TCB extension */
        OS_STK        *OSTCBStkBottom;    /* Pointer to bottom of stack                */
        INT32U         OSTCBStkSize;      /* Size of task stack (in number of stack elements) */
        INT16U         OSTCBOpt;          /* Task options as passed by OSTaskCreateExt() */
        INT16U         OSTCBId;           /* Task ID (0..65535)                        */
    #endif

    struct os_tcb *OSTCBNext;              /* Pointer to next TCB in the TCB list      */
    struct os_tcb *OSTCBPrev;              /* Pointer to previous TCB in the TCB list   */

    #if ((OS_Q_EN > 0) && (OS_MAX_QS > 0)) || (OS_MBOX_EN > 0) || (OS_SEM_EN > 0) || (OS_MUTEX_EN > 0)
        OS_EVENT      *OSTCBEvtPtr;        /* Pointer to event control block            */
    #endif
}

```

TCB

```
#if ((OS_Q_EN > 0) && (OS_MAX_QS > 0)) || (OS_MBOX_EN > 0)
    void          *OSTCBMsg;          /* Message received from OSMBboxPost() or OSQPost() */
#endif

#if (OS_VERSION >= 251) && (OS_FLAG_EN > 0) && (OS_TASK_EN > 0)
    OS_FLAG_NODE *OSTCBFlagNode;      /* Pointer to event flag node */
#endif

    OS_FLAGS      OSTCBFlagsRdy;      /* Event flags that made task ready to run */

    INT16U        OSTCBDly;            /* Nbr ticks to delay task or, timeout waiting for event */
    INT8U         OSTCBStat;           /* Task status */
    INT8U         OSTCBPrio;           /* Task priority (0 == highest, 63 == lowest) */

    INT8U         OSTCBX;              /* Bit position in group corresponding to task priority (0..7) */
    INT8U         OSTCBY;              /* Index into ready table corresponding to task priority */
    INT8U         OSTCBBitX;           /* Bit mask to access bit position in ready table */
    INT8U         OSTCBBitY;           /* Bit mask to access bit position in ready group */

    BOOLEAN       OSTCBDelReq;         /* Indicates whether a task needs to delete itself */
#endif
} OS_TCB;
```

Utilisé si l'option OSTCBOpt.OS_TASK_EN = 1

wait(timeout) ou wait(event, timeout)

Évite les calculs en-ligne, cf. chapitre suivant

OSTaskDel


- `OSTaskDel(OS_PRIO_SELF);`
- Auto Destruction of a task when it has finished processing so that it does not take up memory unnecessarily.

EXECUTION AND SCHEDULING OF TASKS ON AN RTOS

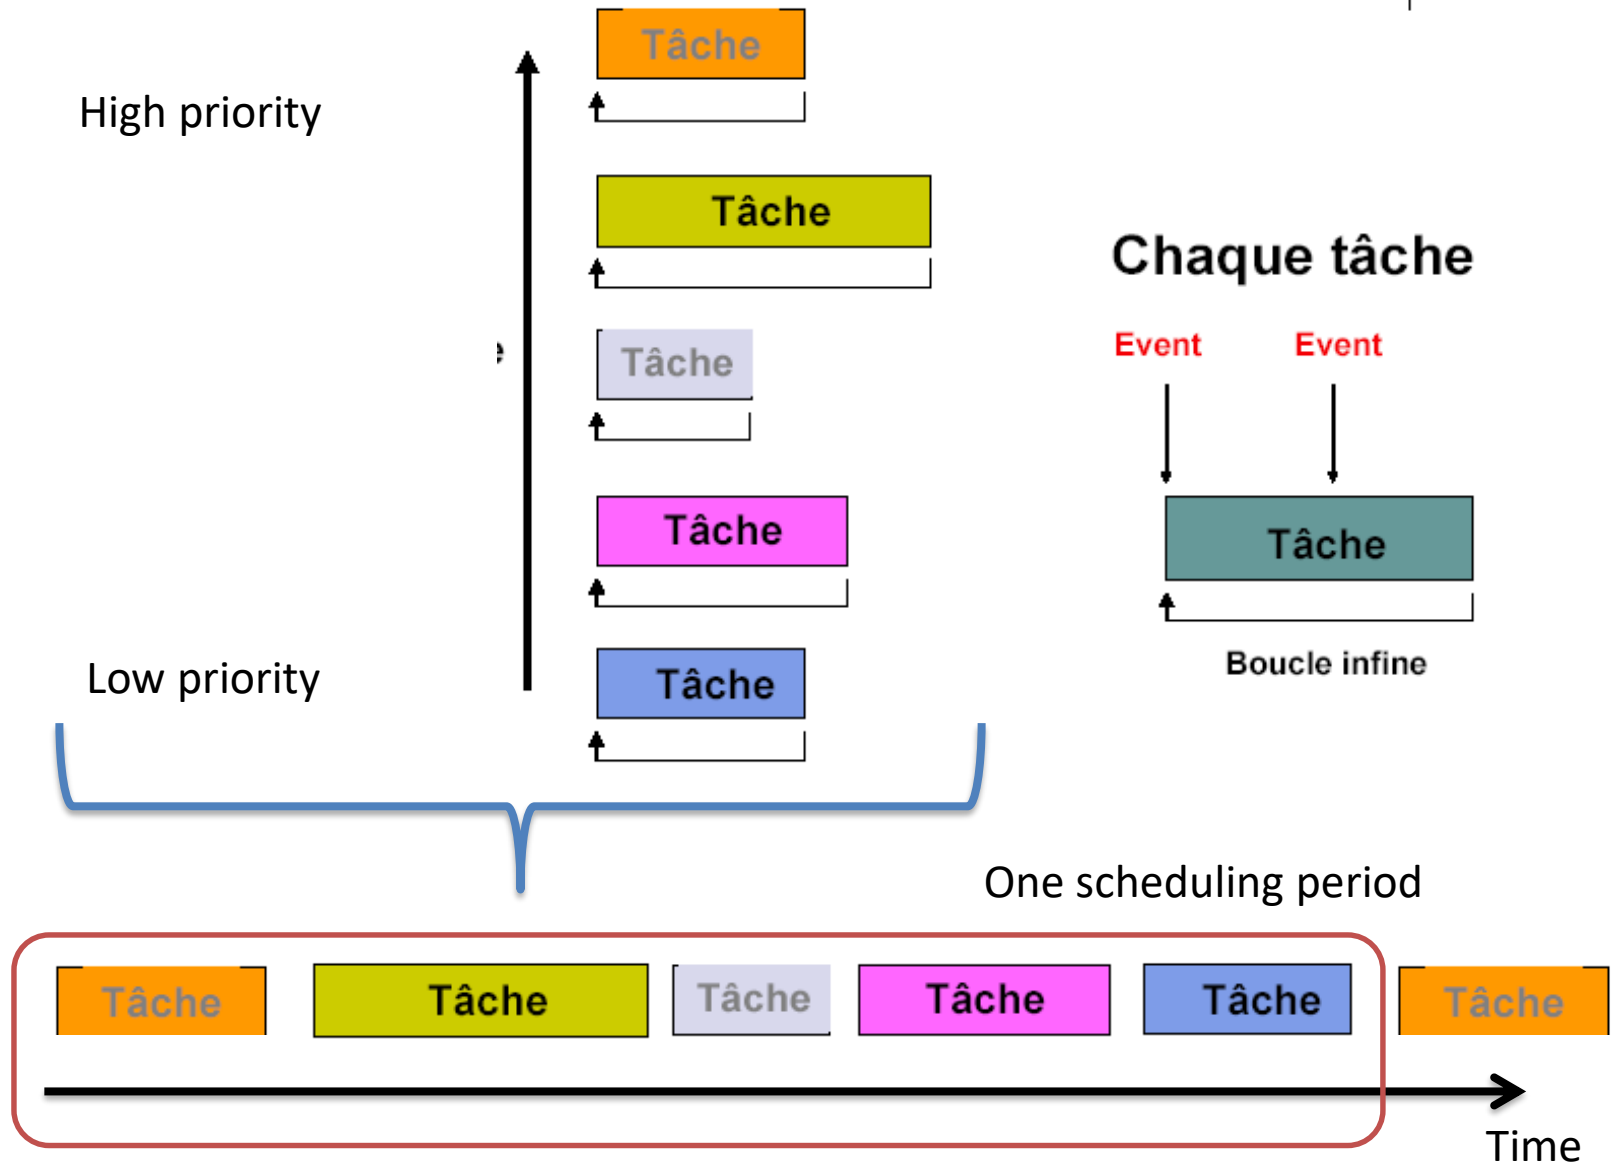
2 types of task

```
Void Run-to-completion_Task (void *pdata){  
    Initialization  
    Create another task  
    Create other kernel objects  
    Self-destruction (end of the task)  
}
```

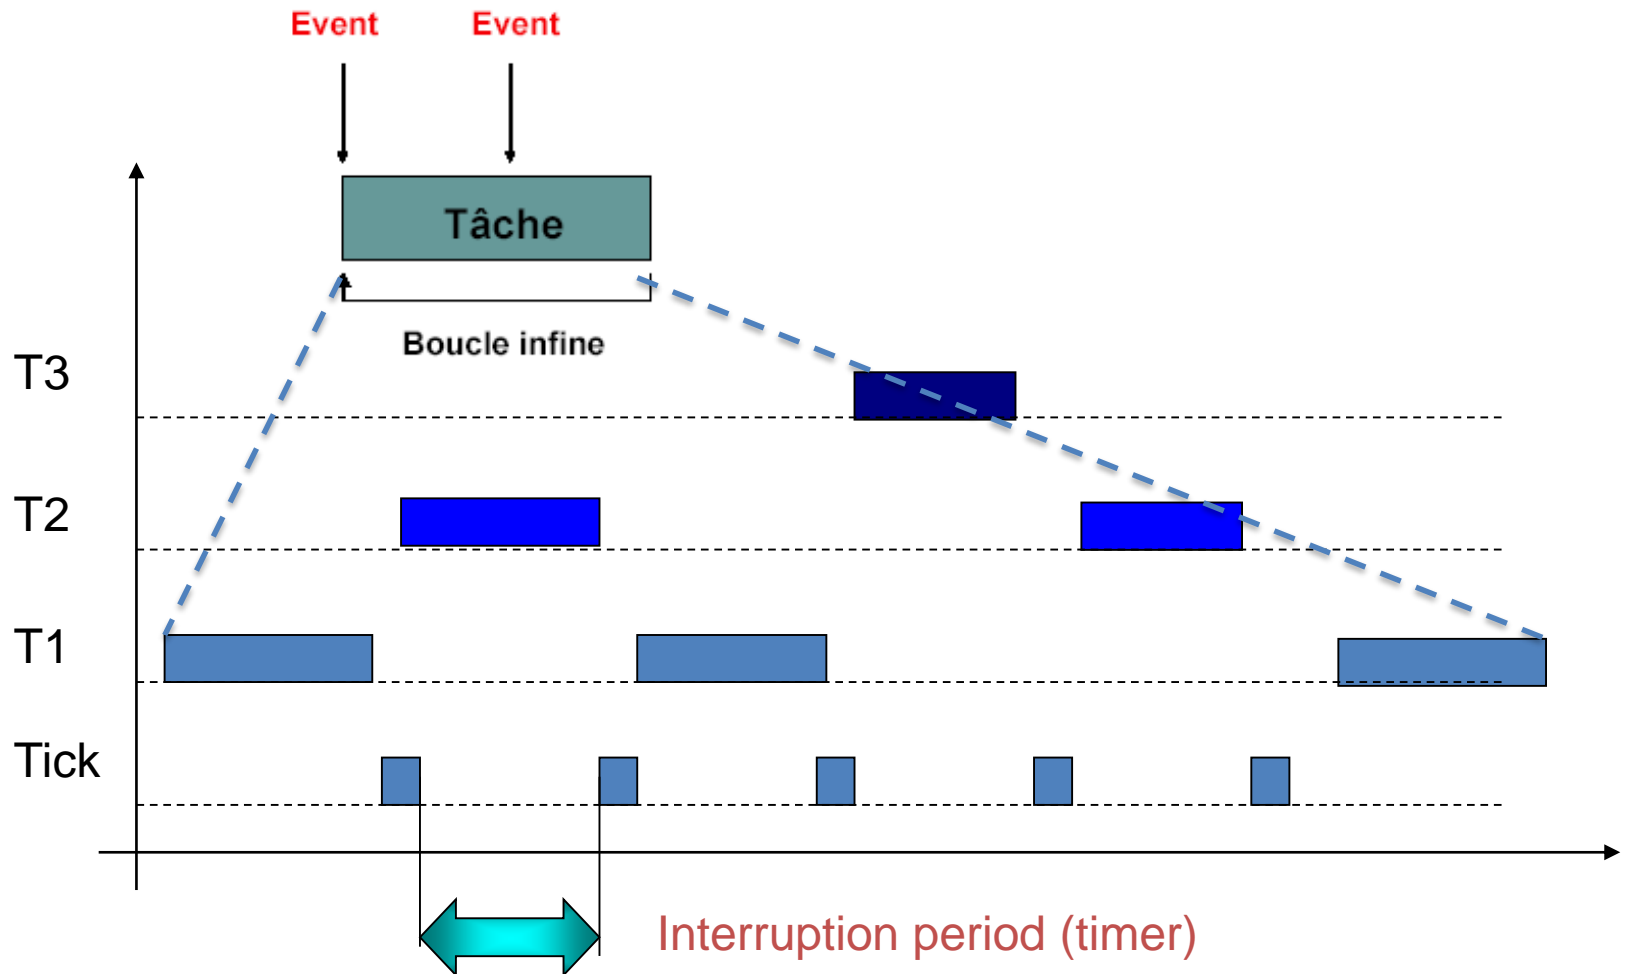
```
Void Endless-loop (void =pdata) {  
    While(1) {  
        Loop body  
        Blocking call  
    }  
}
```



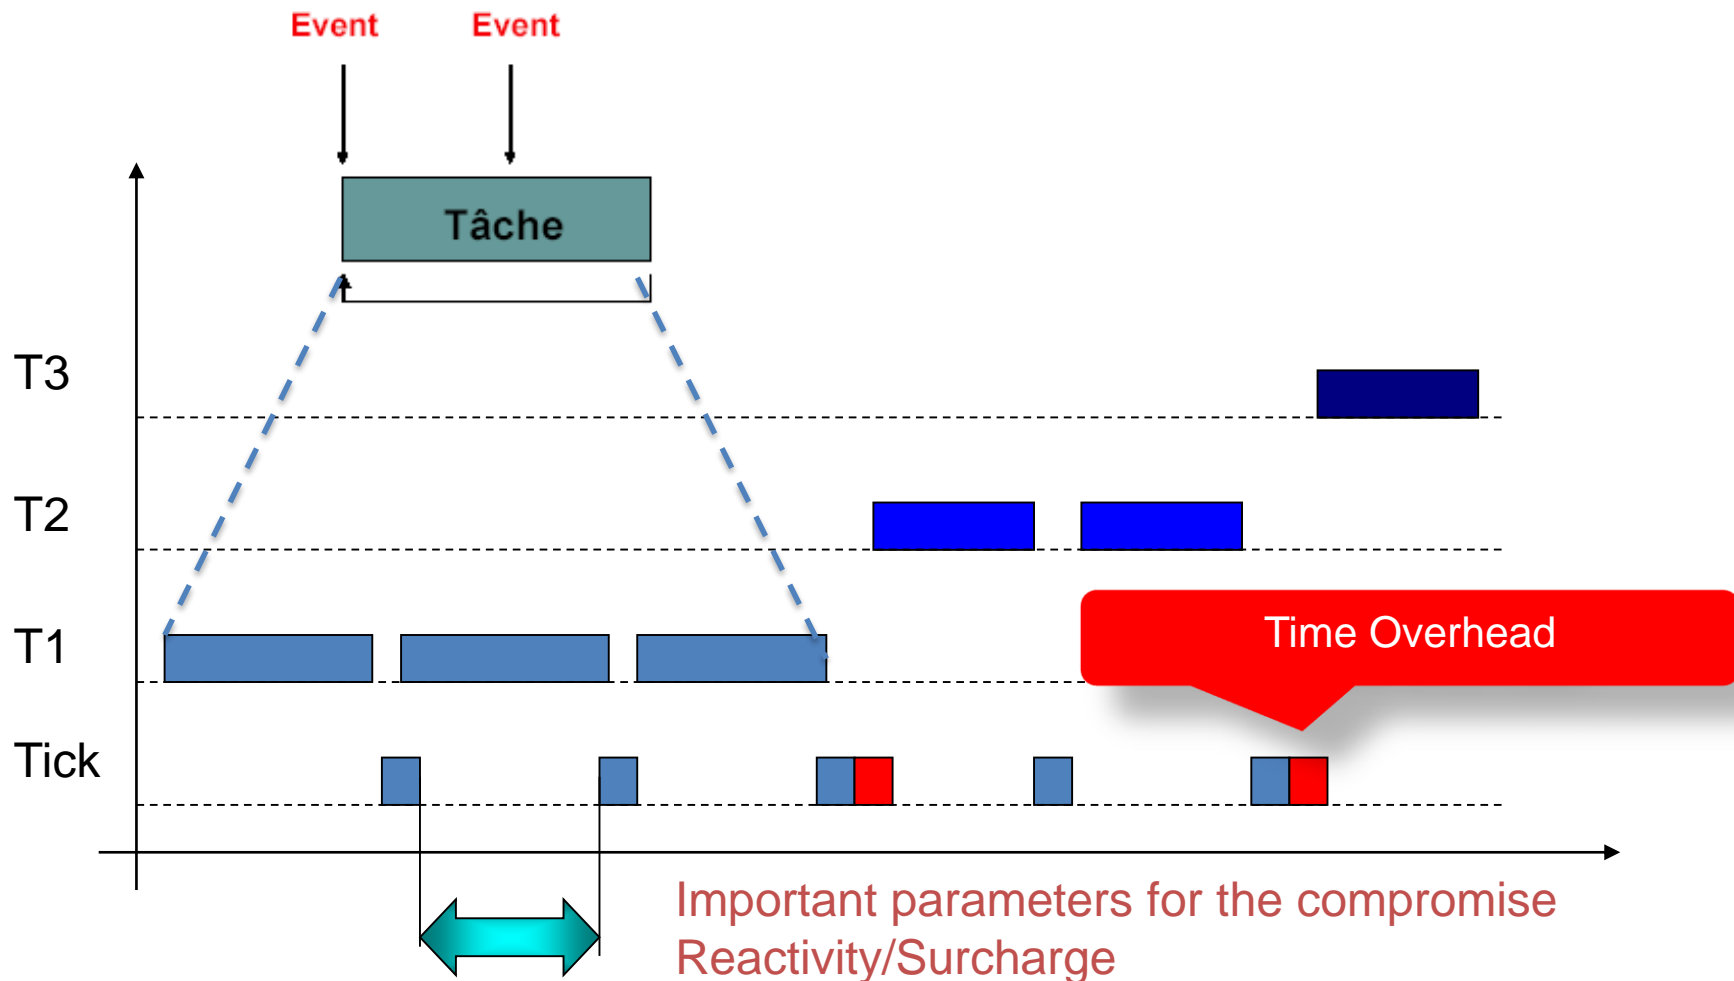
Priority and scheduling



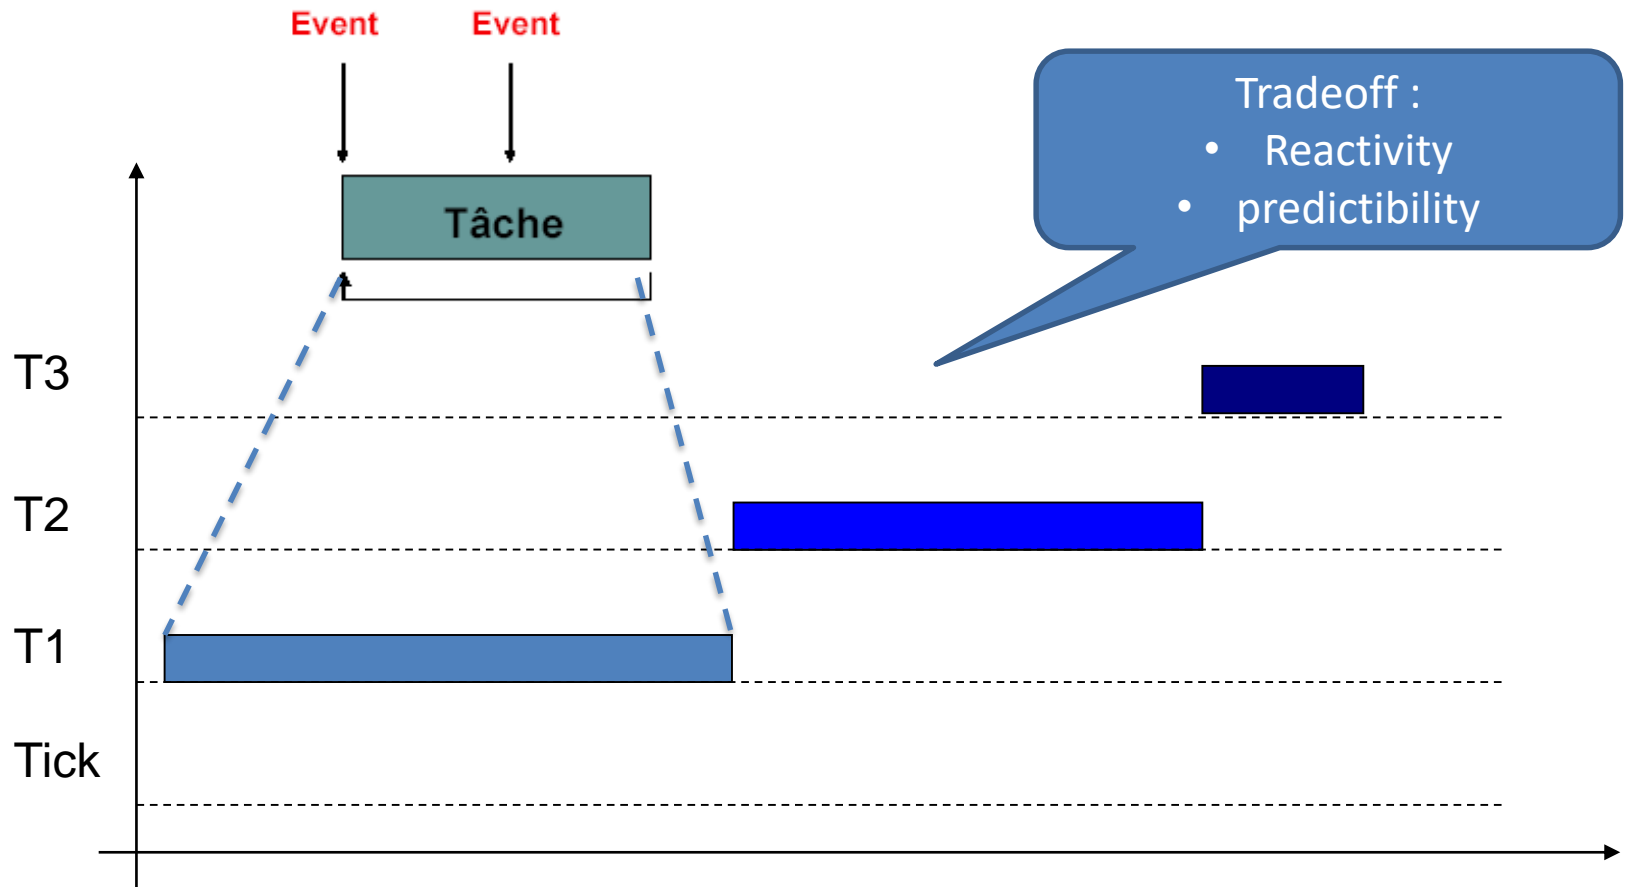
Preemptive OS



Preemptive OS with fixed-priorities



Non-preemptive OS



Reentry functions

Any function called by a task
is executed in the context of the task !


Two types of implementations are possible:

- Reentry functions :

```
void strcpy(char *dest, char *src){  
    while(*dest++ = *src++);  
    *dest = NULL;  
}
```

- Non-reentry functions:

```
int temp;  
void swap(int *x, int *y){  
    temp = *x;  
    *x = *y;  
    *y = temp;  
}
```



Volatile and static keywords

- **Volatile**

- Used in the case of variables whose value can change spontaneously:
 - without processor action, memory mapped devices
 - By another task in case of multithreaded software
- This prefix tells the compiler to avoid optimizations that generate a systematic memory read instruction.
 - More often used in embedded programming

- **Static**

- On a variable, keeps the value of the local variable (allocation out of stack)
- On a function limits the definition of the symbol inside the object (file)

Volatile, example 3:

Preemption and shared variables

- The compiler cannot have knowledge of context switches
- Whereas the shared variables can be changed at any time by another function

```
int cntr;  
void task1(){  
    cntr = 0;  
    while(cntr == 0){  
        // do something  
        sleep(1);  
    }  
    // do something else  
}
```

```
void Task2(){  
    ...  
    cntr++;  
    sleep(10);  
    ...  
}
```

That's why it's safer to declare shared variables as **volatile**

Simple types

- Integers depend on the target architecture, so they are redefined in the specific part: in OS_CPU.h
- On the Nios MCU,
 - typedef unsigned char BOOLEAN; /* Unsigned 8 bit quantity */
 - typedef unsigned char INT8U; /* Signed 8 bit quantity*/
 - typedef signed char INT8S; /* Unsigned 16 bit quantity */
 - typedef unsigned int INT16U; /* Signed 16 bit quantity*/
 - typedef signed int INT16S; /* Unsigned 32 bit quantity */
 - typedef unsigned long INT32U; /* Signed 32 bit quantity */
 - typedef signed long INT32S; /* Single precision floating point */
 - typedef float FP32; /* Double precision floating point */
 - typedef double FP64;

Delays in a RTOS

- uCOS works in real time thanks to the notion of time.
- This time is given by a source called Clock Tick which is a periodic interruption (ISR).
- The period depends on the application.
- Rule 3: ISR = overhead with respect to the system
 - too small and priority tasks will wait longer
 - too big and the extra cost will become a handicap

OSTimeDly()

- A task can itself be put on hold / waiting state
- Calling this function causes a context shift to the next priority task
- Parameter = number of tick between 0 and 65.535
- The task will only be executed again after the time has elapsed and only if it has the highest priority.
- OSTimeDlyHMSM() does the same by taking Hours, Minutes, Seconds and Milliseconds as parameters.
- Rule 4 : Never call this function after disabling SRI!!
- What is the difference with usleep function ?

OSStart()

- It initializes variables and data structures.
- It creates the Idle() task that is always ready to run and has the lowest priority (OS_LOWEST_PRIO = 63).
- If the OS_TASK_STAT_EN and OS_TASK_CREATE_EXT tags are set to 1, it also creates a statistics task with OS_LOWEST_PRIO -1 priority.

OSStart()

- Launches multi-tasking management
- Rule 1 (reminder): You must therefore have already created at least one task
- OSStart finds the TCB of the HPT
- It calls OSStartHighRdy() which is described in OS_CPU.asm depending on the target processor.

Summary

General scheme of a program under uCOS

```
// 1 – Static allocation of memory stacks  
OS_STK TaskStartStk[TASK_STK_SIZE];
```

```
// 2 – declaration of synchro/com objects  
OS_EVENT *mbox, mq, mutex;
```

```
void main (void){
```

```
// 3 - Create at least one task  
OSTaskCreate( TaskStart,  
              (void)*0,  
              &TaskStartStk[TASK_STK_SIZE -1],  
              TASK_START_PRIO);  
OSStart(); // start multi-tasking  
}
```

```
void TaskStart(void *pdata){  
    // 4 – Create other tasks  
}
```

Inter-Process Communication :

IPC

Communication between tasks

i) Shared Ressources

- The easiest way to make 2 tasks communicate is to use a shared memory area.
- Especially when they run in a single address space.
- However, this requires ensuring exclusive access by each task at a given time.
- This translates into several methods:
 - a) Stop interrupts, // CRITICAL SECTION
 - b) Test-And-Set operation
 - c) Stop the scheduler,
 - d) Use a semaphore,

a) Critical Sections

- To prevent other tasks or ISRs from modifying a critical section, it is necessary to interrupt the interruptions.
- interrupt downtime = important parameter of an RTOS (interrupt latency)
- Depends on the processor and the application
- In uCOS, 2 macros are used from OS_CPU.H :
 - `OS_ENTER_CRITICAL()`
 - `// critical section`
 - `OS_EXIT_CRITICAL()`

Interrupt latency

- One of the most important characteristics of an RTOS = the time during which interruptions are stopped (critical sections).
- IRQs are used in an on-board environment as a means of initiating user code triggered by an external asynchronous sensor (ABS braking, sensor sample...).

Interruptions

- Interruptions are asynchronous events triggered by hardware mechanisms.
- When the CPU receives an interrupt, it saves the context of the current task and connects to the routine corresponding to the IRQ number in its interrupt vector.
- At the end of the routine, the CPU returns to :
 - The highest priority task (preemptive mode)
 - The interrupted task (non-preemptive mode)

b) Operation TAS

- This operation is used when the system does not have a RT kernel.
- Principle test and set in a single clock cycle:
 - Test of the value of a global variable
 - If Val = 0
 - The function has access to the resource
 - It sets the variable to 1 // Test and Set
- Some processors implement this service in hardware in their instruction set.

c) Stop the scheduler

- Rough solution that does not interrupt interruptions, only multi-tasking.
- For example: :
 - OSSchedLock()
 - Access to share data
 - OSSchedUnlock()

d) Semaphores

- Invented in the 1960s by Edgser Dijkstra
- Binary semaphores // Mutex
- Counting semaphores
- Three types of operations :
 - Create (initialize)
 - Wait (Pend) ($\text{sem} > 0$)? $\text{sem}--$: `wait()`;
 - Signal (Post) ($\text{sem} == 0$)? $\text{sem}++$: `notify()`;
- The task that launches (notify) is either
 - Highest priority (uC/OS)
 - The one who requested it first (FIFO)

Assessment of shared resource access

- For the access to a shared area, the semaphore is the least risky solution.
- If other solutions are misused, the consequences can be much more serious.
- However, for a simple access to a 32-bit variable, stopping interrupts will be less expensive than using a semaphore without changing the interrupt latency!

Deadlock

- Deadlock occurs when 2 tasks are waiting for access to resources taken by the other one (see course on scheduling).
- To avoid these situations the tasks must :
 - Acquire resources in the same order,
 - Release resources in reverse order.
 - RTOS kernels also allow you to specify a **Timeout** on waiting for semaphores.
 - These calls return a different error code to inform the task that the wait was not successful.

ii) Sending messages

- Message **Mailbox**
 - Sending a pointer to a mailbox
 - Only one letter at a time !
- Message **Queue**
 - Sending several messages (mailbox queue[] ;)
 - Reading in FIFO mode
 - A list of pending processes is built by the scheduler.

State diagram under ucOS

