

BITTORRENT

BitTorrent

2

- The technology has three aspects
 - ▣ Many BitTorrent clients are available in open-source
 - Protocol initially written in English (no math, no pseudocode)
 - Attempt of standardize (see bittorrent.org)
 - ▣ Many existing clients, mostly open-source
 - ▣ A clever idea: using “tit-for-tat” mechanisms to reward good behavior and to punish bad behavior
- This third aspect is especially intriguing!

The basic BitTorrent Scenario

3

- Millions want to download the same popular huge files (for free)
 - ▣ ISO's
 - ▣ Media files (the majority !)
- Client-server model fails
 - ▣ Single server fails
 - ▣ Can't afford to deploy enough servers


Why not use IP Multicast?

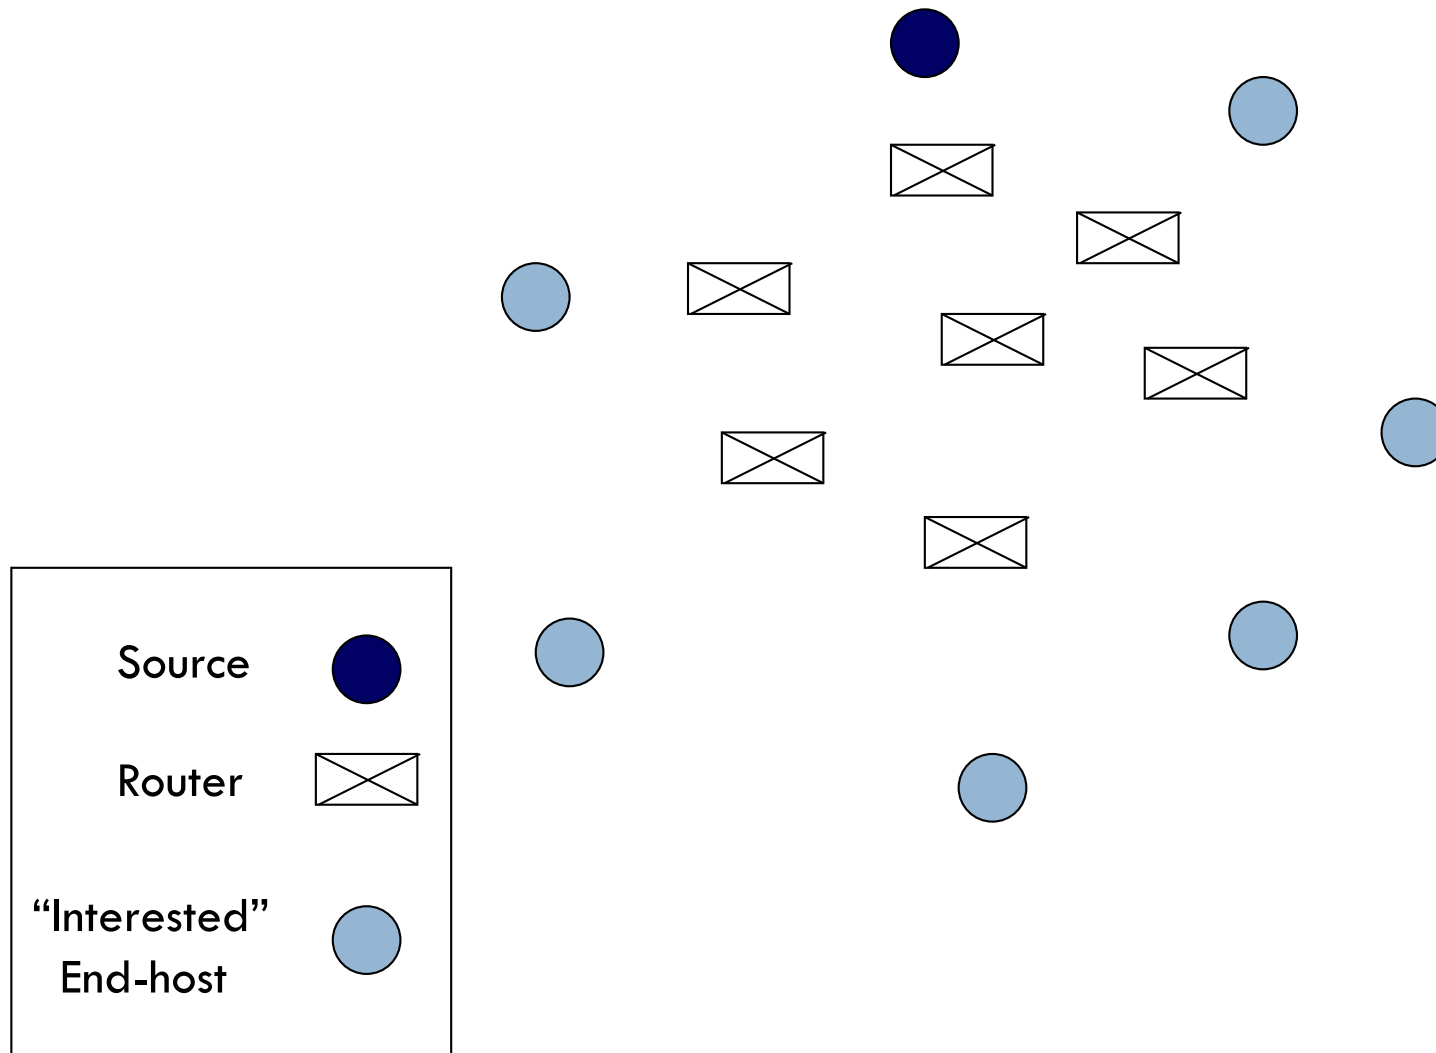
4

- IP Multicast not a real option in general WAN settings
 - ▣ Not supported by many ISPs
 - ▣ Most commonly seen in private data centers
- Alternatives
 - ▣ End-host based Multicast
 - ▣ BitTorrent
 - ▣ Other P2P file-sharing schemes

IP multicast (only with IPv6)

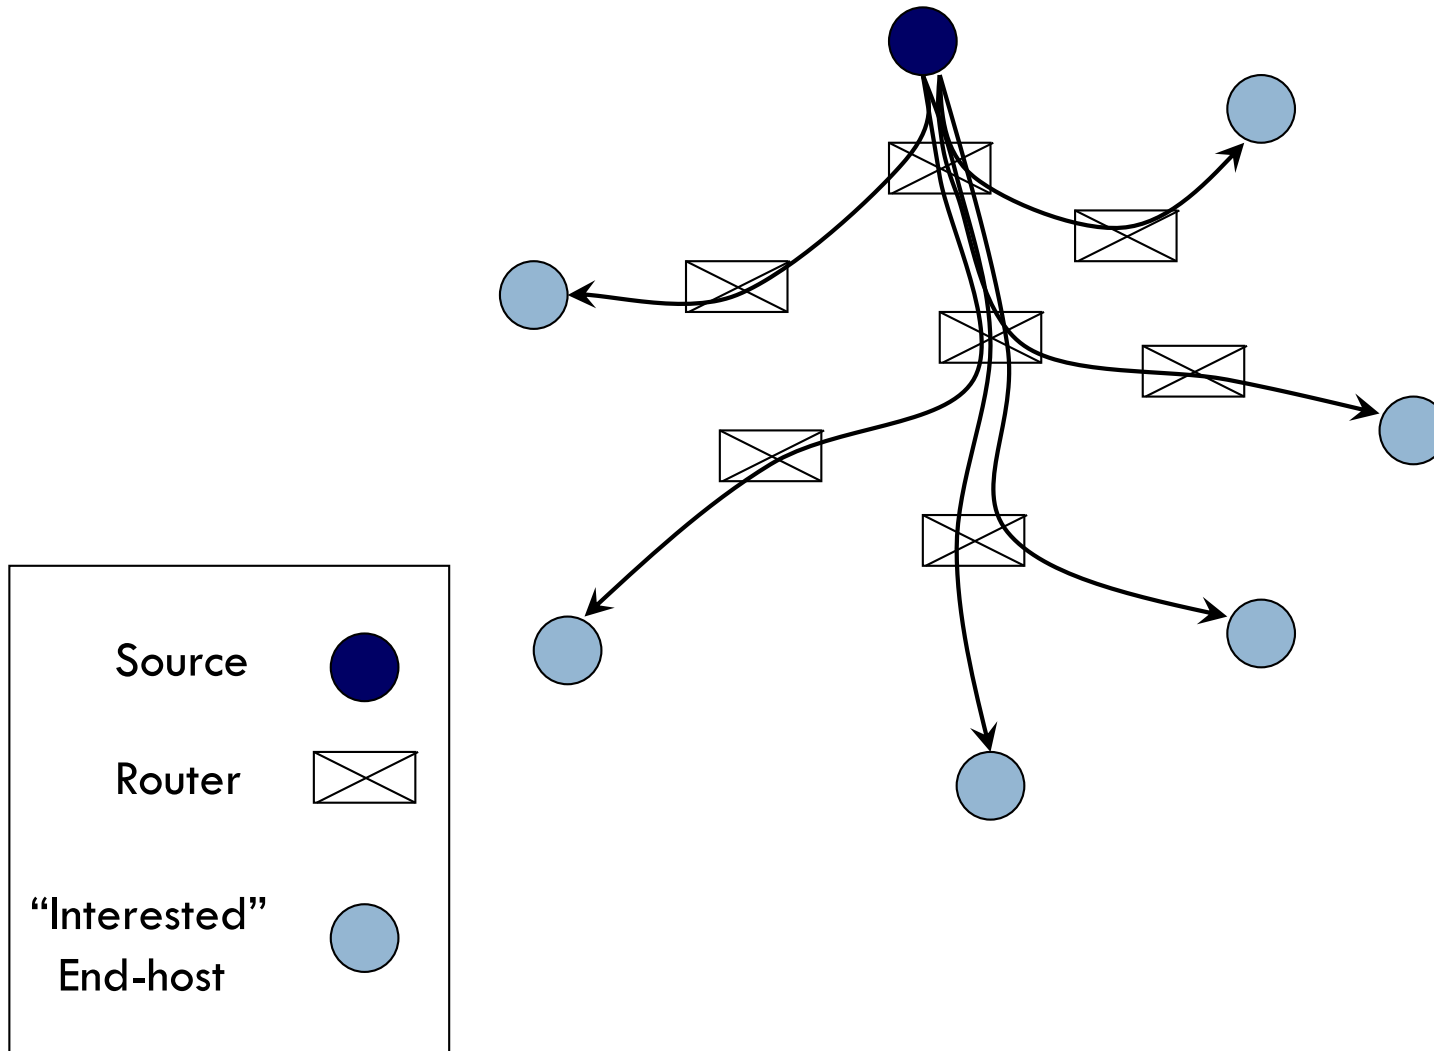
5






Client-Server

7





Client-Server

8



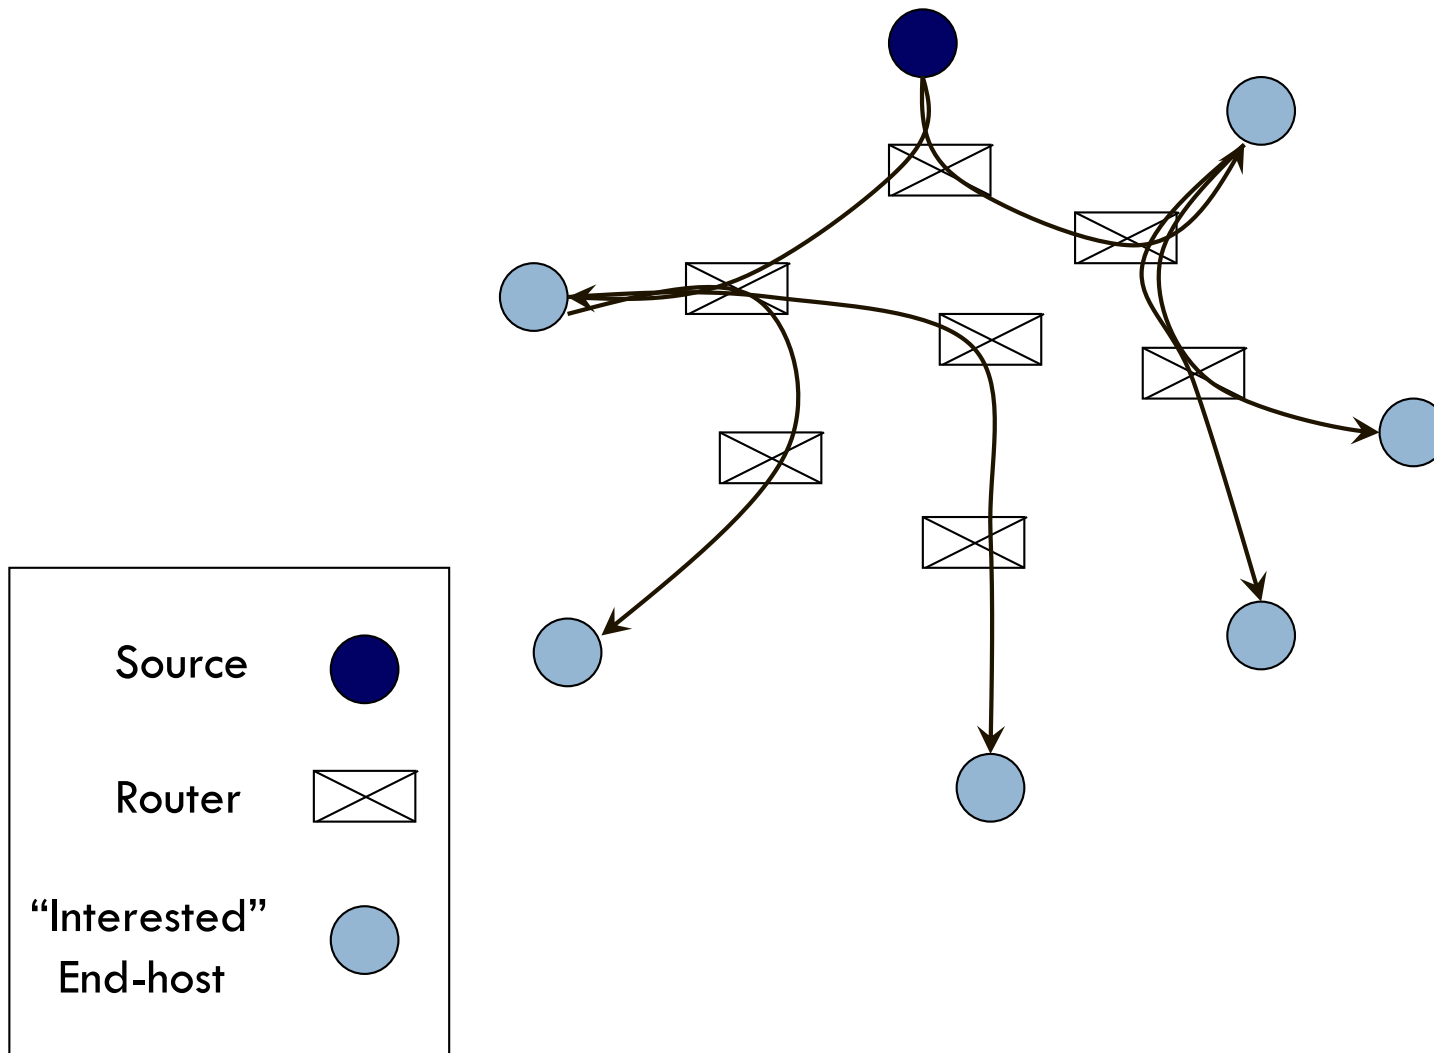
(plain) IP multicast

9



End-host based multicast

10



End-host based multicast

11

- **Single-uploader** versus **Multiple-uploaders**
 - ▣ Lots of nodes want to download
 - ▣ Make use of their uploading abilities as well
 - ▣ Node that has downloaded (part of) file will then upload it to other nodes
 - Uploading costs amortized across all nodes


End-host based multicast

12

- Also called **Application-level Multicast**
- Many protocols proposed this two decades ago !
- All use single trees
 - Problem with single trees ? Unfortunately yes !


End-host multicast using single tree

13



End-host multicast using single tree

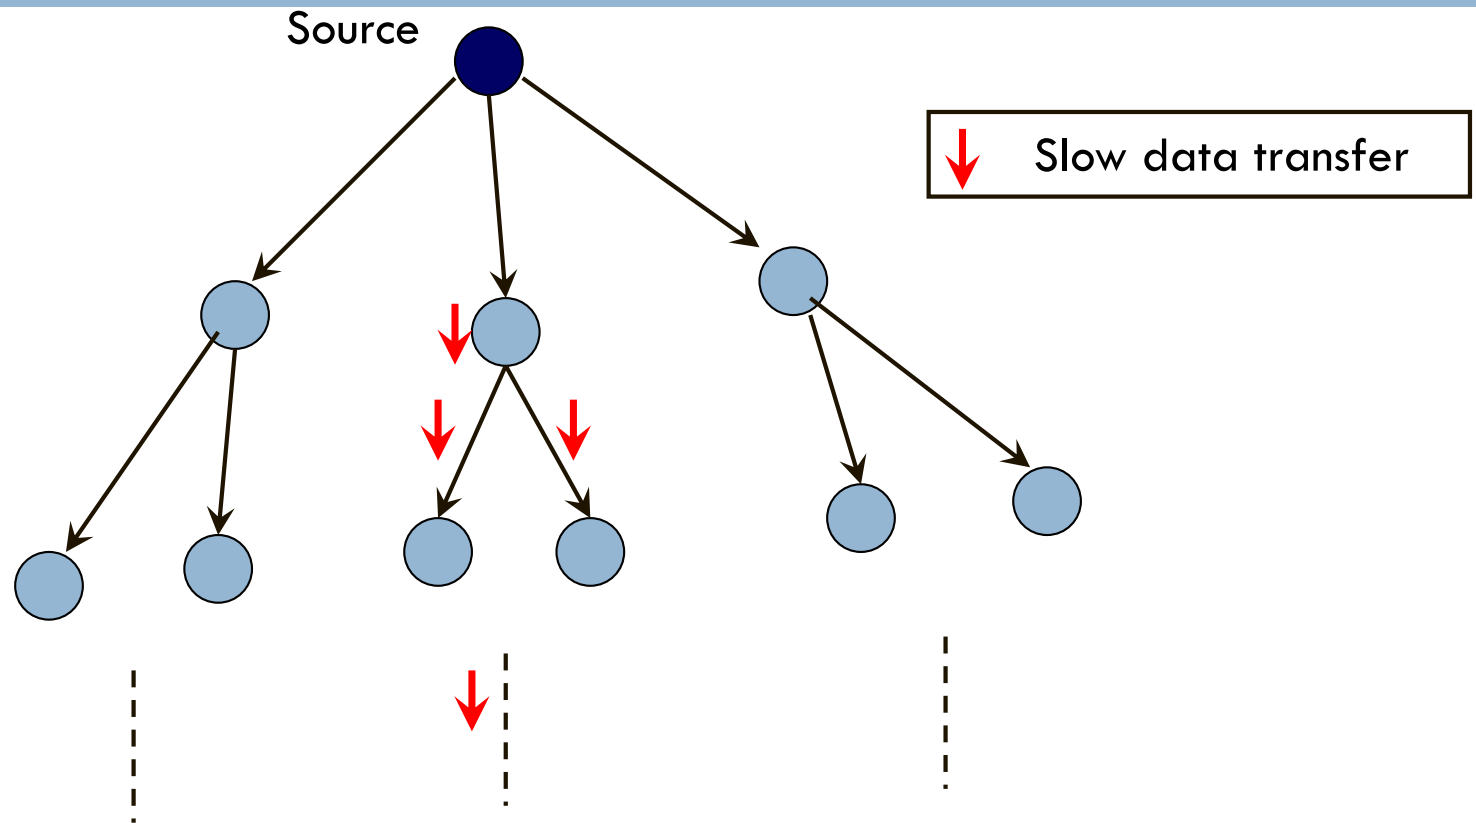
14



A supernode failure
implies that all
subtree nodes fails

End-host multicast using single tree

15




End-host multicast using single tree

16

- Tree structure is **push-based** – node receives data, pushes data to children
- Failure of **interior-nodes** affects downloads in entire subtree rooted at node
- Slow interior-nodes similarly affects entire subtree
- Also, leaf-nodes don't do any sending!
- Though later multi-tree / multi-path protocols mitigate some of these issues, **tree are not a good topology !!!**

BitTorrent

Give and ye shall receive!



17

- Written by Bram Cohen (in Python) in 2001
- “Pull-based” & “Swarming” approach
 - ▣ Each file split into smaller pieces
 - ▣ Nodes request desired pieces from neighbors
 - As opposed to parents pushing data that they receive
 - ▣ Pieces not downloaded in sequential order
 - ▣ Previous multicast schemes aimed to support “streaming”; BitTorrent does not (but new bittorrent dialects supports streaming ... e.g. <201X Spotify used P2P tech)
- Encourages contribution by all nodes

in 2018

EMEA ↑ APPLICATION TRAFFIC SHARE **TOP 10**


1 BITTORRENT TRANSFER
31.73% ↑


2 GOOGLE
9.42% ↑


18


in 2010


Internet Traffic




































BitTorrent Swarm

29

- **Swarm** (Essaim/Enjambre/Sciame/埃賽姆)
 - ▣ Set of peers all downloading the same file
 - ▣ Organized in a **random mesh** topology
- Each node knows list of pieces downloaded by neighbors
- First intuition: a node requests pieces it does not own from neighbors

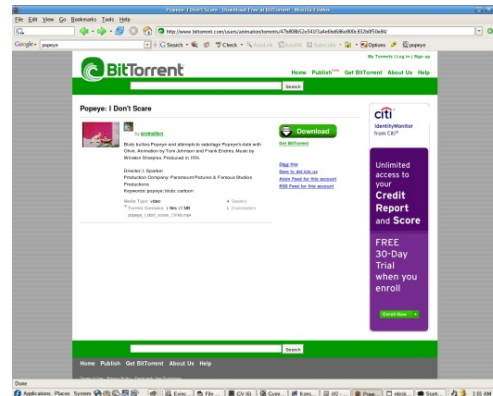

How a node enters a swarm for file “**popeye.mp4.torrent**”



- File **popeye.mp4.torrent** hosted at a (well-known) webserver
- The .torrent has the address of the **TRACKER** for the file
- The tracker, which runs on a webserver as well, keeps track of all peers downloading the file. It refresh the track list continuously (approx 30s)

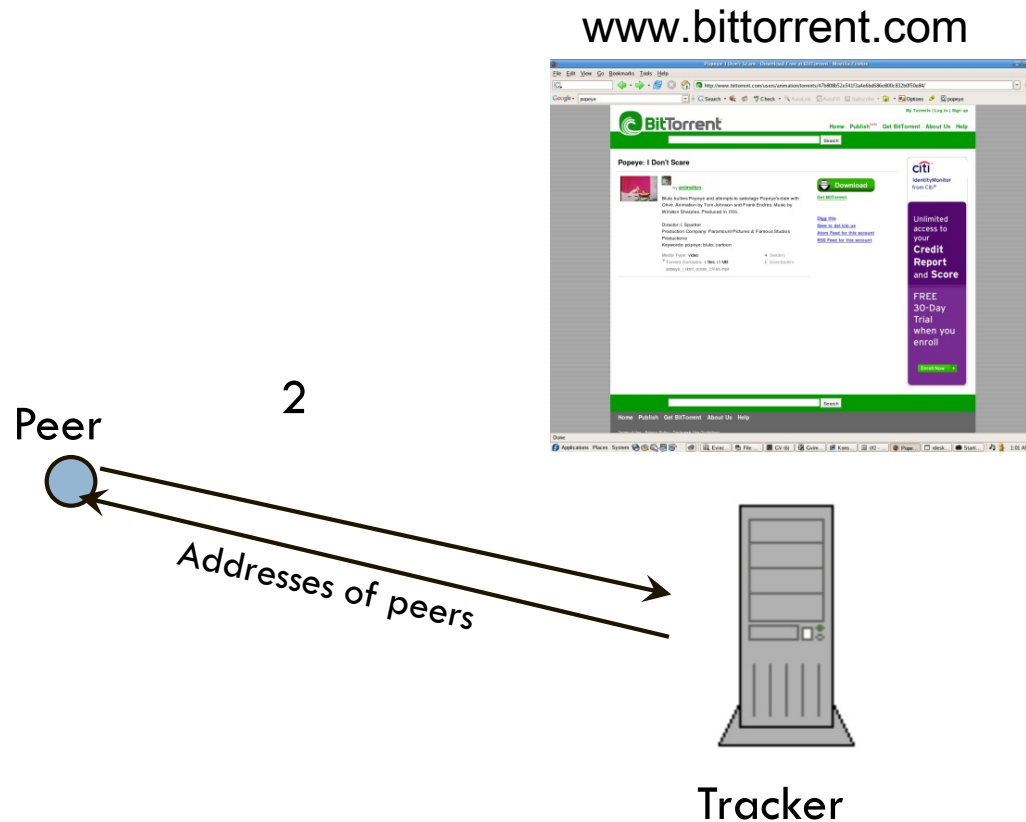
How a node enters a swarm for file “madonnaCD.mp4”

e.g. piratebay



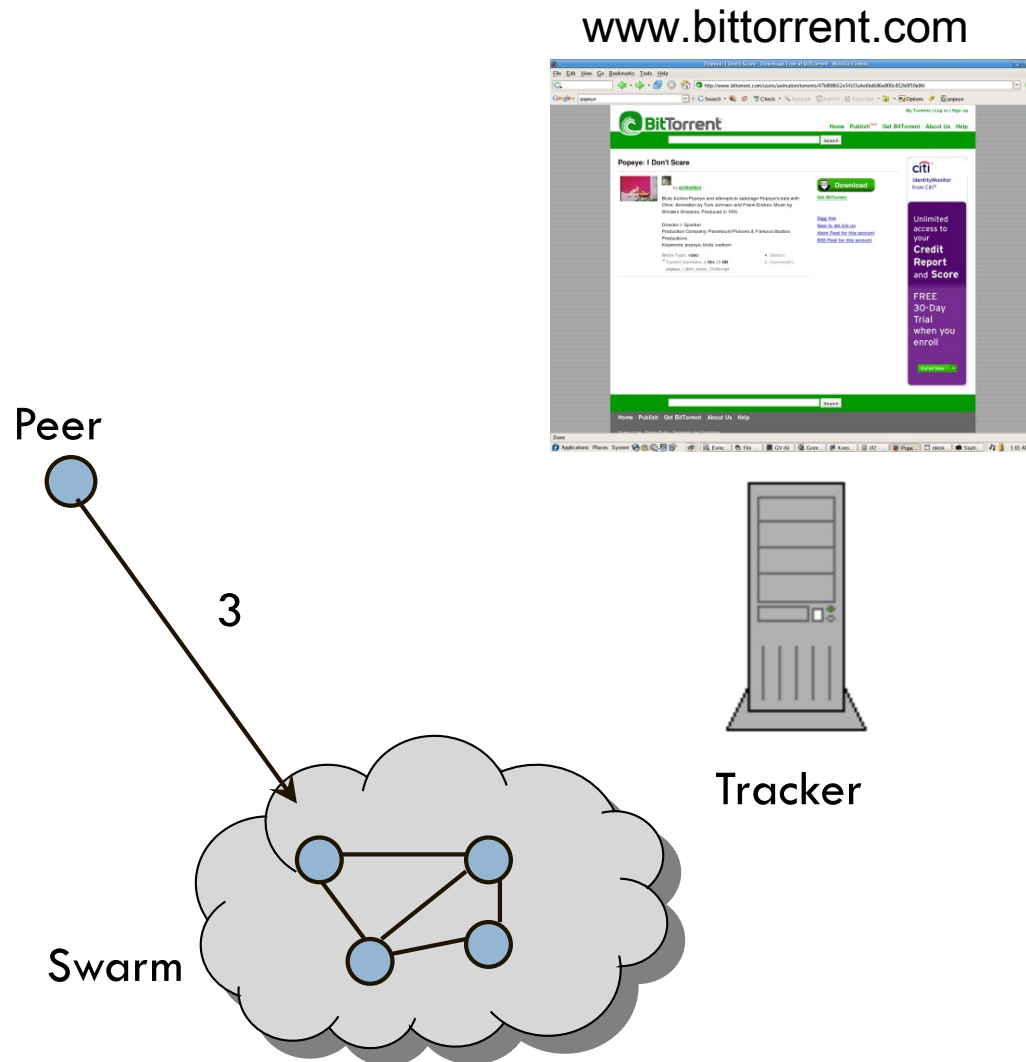
- File **popeye.mp4.torrent** hosted at a (well-known) webserver
- The .torrent has the address of the **TRACKER** for the file
- The tracker, which runs on a webserver as well, is a **SERVER** that keeps track of all peers downloading the file. It refresh the track list continuously (approx 30s)

How a node enters a swarm for file “popeye.mp4”



- File **popeye.mp4.torrent** hosted at a (well-known) webserver
- The .torrent has the address of the **TRACKER** for the file
- The tracker, which runs on a webserver as well, keeps track of all peers downloading the file. It refresh the track peer list continuously (approx 30s)

How a node enters a swarm for file “popeye.mp4”



- File **popeye.mp4.torrent** hosted at a (well-known) webserver
- The .torrent has address of **tracker** for file
- The tracker, which runs on a webserver as well, keeps track of all peers downloading the file. It refresh the track list continuously (approx 30s)

Contents of .torrent file

34

- HTTP/URL/UDP of tracker
- Piece length – Usually 256 KB up to 2 MB
- SHA-1 hashes of each piece in file
 - ▣ For reliability/data integrity
- **Files names and file hierarchy** – allows download of multiple files, e.g. a directory containing files

Inside a .torrent file

- The file is encoded using an original *B-encoding*
 - ▣ Announce URL/HTTP/UDP of the tracker
 - ▣ Some optional fields
 - Creation date, comment, created by
- *Info key*
 - ▣ Length on the content in bytes
 - ▣ md5 hash of the content
 - ▣ Pieces SHA-1 hash are enough
 - ▣ File Name
 - ▣ Piece length (256kB, 512kB, 1024kB, etc.)
 - ▣ Concatenation of all pieces SHA-1 hash

BENCODING (in a torrent file)

36

bencoding


- Strings are length-prefixed base ten followed by a colon and the string. For example `4:spam` corresponds to 'spam'.
- Integers are represented by an 'i' followed by the number in base 10 followed by an 'e'. For example `i3e` corresponds to 3 and `i-3e` corresponds to -3. Integers have no size limitation. `i-0e` is invalid. All encodings with a leading zero, such as `i03e`, are invalid, other than `i0e`, which of course corresponds to 0.
- Lists are encoded as an 'l' followed by their elements (also bencoded) followed by an 'e'. For example `l4:spam4:eggse` corresponds to ['spam', 'eggs'].
- Dictionaries are encoded as a 'd' followed by a list of alternating keys and their corresponding values followed by an 'e'. For example, `d3:cow3:moo4:spam4:eggse` corresponds to {'cow': 'moo', 'spam': 'eggs'} and `d4:spam11:a1:bee` corresponds to {'spam': ['a', 'b']}. Keys must be strings and appear in sorted order (sorted as raw strings, not alphanumerics).

Pieces and Blocks

37

Pieces and Blocks

- Content is split into *pieces* (256KB-2MB)
- Pieces are split into *blocks* (16KB)



Terminology

38

- **Seed**: peer with the entire file
 - ▣ Original Seed: the first seed
- **Leech**: peer that's downloading the file
 - ▣ Fairer term might have been “downloader”
- **Block**: Further subdivision of a piece
 - ▣ The “unit for requests” is a block (16KB)
 - ▣ A peer uploads only **after** assembling complete piece

Peer-peer transactions:

Choosing pieces to request 1

39

- **Rarest-first:** Look at all pieces at all peers, and request piece that's owned by fewest peers
 - ▣ Increases diversity in the pieces downloaded
 - avoids case where a node and each of its peers have exactly the same pieces; increases throughput
 - ▣ Increases likelihood all pieces still available even if original seed leaves before any one node has downloaded entire file

Choosing pieces to request 2

40

□ Random First Piece:

- When peer starts to download, request random piece.
 - So as to assemble first complete piece quickly
 - Then participate in uploads
- When first complete piece assembled, switch to rarest-first

Choosing pieces to request 3

41

□ End-game mode:

- When requests sent for all blocks, (re)send requests to all peers
- To speed up completion of download
- ``Cancel`` request for downloaded blocks

Tit-for-tat as incentive to upload

42

- Want to encourage all peers to contribute
- Peer *A* said to **choke** peer *B* if it (*A*) decides not to upload to *B*
- Each peer unchokes at most 4 *interested* peers at any time
 - ▣ The three with the largest upload rates to *A*
 - Where the tit-for-tat comes in
 - ▣ Another randomly chosen (**Optimistic Unchoke**)
 - To periodically look for better choices

Tit-for-tat

43

- Best deterministic strategy for the ``Prisoner's Dilemma``
- 1 - Unless provoked, the agent will always cooperate
- 2 - If provoked, the agent will retaliate
- 3 - The agent is quick to forgive
- The agent must have a good chance of competing against the opponent more than once

Anti-snubbing

44

- A peer is said to be **snubbed** if each of its peers chokes it
- To handle this, snubbed peer stops uploading to its peers
- Optimistic unchoking done more often
 - ▣ Hope is that will discover a new peer that will upload to us

Why BitTorrent took off 1

45

- Better performance through “pull-based” transfer
 - ▣ Slow nodes don’t bog down other nodes
- Allows uploading from hosts that have downloaded parts of a file
 - ▣ In common with other end-host based multicast schemes

Why BitTorrent took off 2

46

- Practical Reasons (perhaps more important!)
 - ▣ Working implementation in Python by Bram Cohen with simple well-defined interfaces for plugging in new content
 - ▣ Many recent competitors got sued / shut down
 - Napster, Kazaa, EMule, ...
 - ▣ Doesn't do “search” per se. Users use well-known, trusted sources to locate content
 - Avoids the “*pollution problem*”, where garbage is passed off as authentic content

Pros and cons of BitTorrent 1

47

□ Pros

- ▣ Proficient in utilizing partially downloaded files
- ▣ Discourages “freeloading”/ “free riders”
 - By rewarding fastest uploaders
- ▣ Encourages diversity through “rarest-first”
 - Extends lifetime of swarm
- Works well for very popular contents

Pros and cons of BitTorrent 2

48

□ Cons

- Assumes all interested peers active at same time; performance deteriorates if swarm “cools off”

Pros and cons of BitTorrent 3

49

- Dependence on centralized tracker: pro/cons ?
 - ▣ ☹ Single point of failure: new nodes can't enter swarm if tracker goes down
 - ▣ Lack of a search feature
 - ☺ Prevents pollution attacks
 - ☹ Users need to resort to out-of-band search: well known torrent-hosting sites / plain classic web-search

“Trackerless” BitTorrent

50

- To be more precise, “BitTorrent without a centralized-tracker”
- E.g.: Vuze
- Uses a Distributed Hash Table (Kademlia DHT) as a “distributed tracker”
- Tracker run by a normal end-host (not a web-server anymore)
 - ▣ The original seeder could itself be the tracker
 - ▣ or have a node in the DHT randomly picked to act as the tracker

Magnet links

51

Linux Mint Debian [201101] [ISO] [32-Bit] [geno7744]

Type:	Applications > UNIX	Quality:	+1 / -0 (+1)
Files:	1	Uploaded:	2011-01-08 06:30:09 GMT
Size:	986.05 MiB (1033945088 Bytes)	By:	geno7744
Tag(s):	linux linux mint mint debian ubuntu debian	Seeders:	23
		Leechers:	1
		Comments:	1

Info Hash:
2E99D97F1768644A86A8E99BFD80C816490F959B

 [Tweet](#)

Download

Enjoy Movies, TV Shows, Music and Games on your browser!

 **GET THIS TORRENT**
(Problems with magnets links are fixed by upgrading your [torrent client!](#))

Linux Mint Debian Edition (LMDE) is a rolling distribution based on Debian Testing. At the moment, it comes as a 32-bit live DVD with a Gnome desktop.

The purpose of LMDE is to look identical to the main edition and to provide the same functionality while using Debian as a base.

Why is (studying) BitTorrent important?

52

- BitTorrent consumes significant amount of internet traffic today
 - ▣ In 2004, BitTorrent accounted for 30% of all internet traffic (Total P2P was 60%), according to CacheLogic
 - ▣ Slightly lower share in 2005 (possibly because of legal action), but still significant
 - ▣ BT always used for legal software (linux iso) distribution too
 - ▣ Recently: legal media downloads (Fox, BBC), streaming

Implementation (FR)

53

- ❑ Les implémentations les plus connues souhaitent échanger des pièces de taille jusqu'à 2Mb ;
- ❑ Elle préfèrent utiliser plutôt protocoles plus ``élastiques`` comme **UDP** ou **HTTP** : grâce à cela, elle fractionnent ultérieurement la pièce en sous-pièces, ces derniers appelés **BLOCKS** de taille fixe de 16 KB.
- ❑ Il y a un consensus quasi complet dans la communauté des clients bittorrent sur la sous-fragmentation d'une **PIECE** de taille VARIABLE entre 256Kb et 2Gb en **BLOCKS** de taille FIXE de 16kb.

Tit-for-tat

54

Description [[modifier](#) | [modifier le code](#)]

En 1974, [Anatol Rapoport](#) grâce à des études théoriques et empiriques (en partie avec A.M. Chammah), confirmées par deux tournois gagnés en 1979², déduit l'idée que la manière la plus « [efficace](#) » de se comporter vis-à-vis d'autrui est de se comporter comme suit :

1. Attitude de [coopération](#): Dans un premier temps, lorsqu'un individu ou un groupe rencontre un autre individu ou groupe, il a tout intérêt à lui proposer une alliance.
2. Attitude de réciprocité: Dans un second temps, en vertu de la règle de réciprocité, il convient de donner à l'autre en fonction de ce que l'on en reçoit. Si l'autre aide, on l'aide en retour ; si l'autre agresse, il faut répondre en l'agressant à son tour, au coup suivant, de la même manière et avec la même intensité.
3. Attitude de [pardon](#): Dans un troisième temps, il faut [pardonner](#) et offrir de nouveau la coopération.