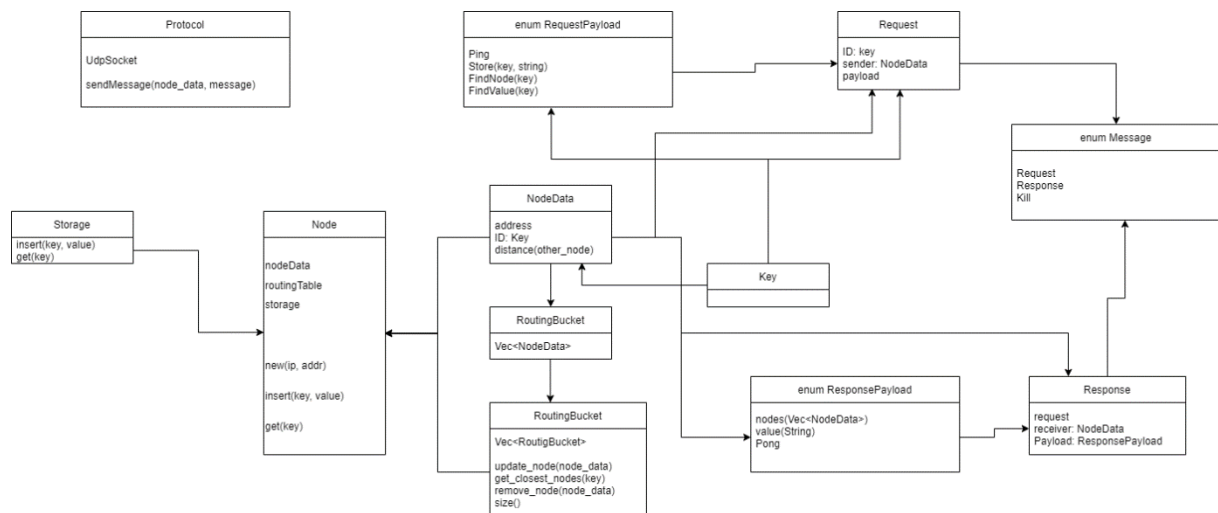


Kademlia code review

<https://gitlab.com/jeffrey-xiao/kademlia-dht-rs>

Architecture :



Messages are serialized with **serde crate** before being sent to the other nodes via UDP.

Comments :

The implementation uses the UDP protocol.

Bucket table :

src/routing.rs

```
97  /// A node's routing table tree.
98  ///
99  /// `RoutingTable` is implemented using a growable vector of `RoutingBucket`. The relaxation of
100 /// k-bucket splitting proposed in Section 4.2 is not implemented.
101 #[derive(Clone, Debug)]
102 pub struct RoutingTable {
103     buckets: Vec<RoutingBucket>,
104     node_data: Arc<NodeData>,
105 }

8   /// A k-bucket in a node's routing table that has a maximum capacity of `REPLICATION_PARAM`.
9   ///
10  /// The nodes in the k-bucket are sorted by the time of the most recent communication with those
11  /// which have been most recently communicated at the end of the list.
12  #[derive(Clone, Debug)]
13  struct RoutingBucket {
14      nodes: Vec<NodeData>,
15      last_update_time: SteadyTime,
16  }
```

The bucket table is a list of list of NodeData. NodeData contains the node ID and address, and allows to sort nodes and get distance between them.

Update bucket table :

All constants are defined in src/lib.rs.

REPLICATION_PARAM is the maximum number of entries in a k-bucket, here 20.

ROUTING_TABLE_SIZE is the maximum number of k-buckets in the routing table, here 256.

```
115     /// Upserts a node into the routing table. It will continue to split the routing table until the
116     /// routing table is full or until the node can be upserted.
117     pub fn update_node(&mut self, node_data: NodeData) -> bool {
118         let distance = self.node_data.id.xor(&node_data.id).leading_zeros();
119         let mut target_bucket = cmp::min(distance, self.buckets.len() - 1);
120
121         if self.buckets[target_bucket].contains(&node_data) {
122             self.buckets[target_bucket].update_node(node_data);
123             return true;
124         }
125
126         loop {
127             /// bucket is not full
128             if self.buckets[target_bucket].size() < REPLICATION_PARAM {
129                 self.buckets[target_bucket].update_node(node_data);
130                 return true;
131             }
132
133             let is_last_bucket = target_bucket == self.buckets.len() - 1;
134             let is_full = self.buckets.len() == ROUTING_TABLE_SIZE;
135
136             /// bucket cannot be split
137             if !is_last_bucket || is_full {
138                 return false;
139             }
140
141             /// split bucket
142             let new_bucket = self.buckets[target_bucket].split(&self.node_data.id, target_bucket);
143             self.buckets.push(new_bucket);
144
145             target_bucket = cmp::min(distance, self.buckets.len() - 1);
146         }
147     }
```

Find_node :

Src/node/mod.rs

```
271     /// Sends a `FIND_NODE` RPC.
272     fn rpc_find_node(&mut self, dest: &NodeData, key: &Key) -> Option<Response> {
273         self.send_request(dest, RequestPayload::FindNode(*key))
274     }
```

It will send a request of type FindNode to the destination, with the requested node ID as parameter.

As a response to a SendNode request, the other node will return the closest nodes to the targeted one:

```

152    /// Handles a request RPC.
153    fn handle_request(&mut self, request: &Request) {
154        info!(
155            "{} - Receiving request from {} {:?}",
156            self.node_data.addr, request.sender.addr, request.payload,
157        );
158        self.clone().update_routing_table(request.sender.clone());
159        let receiver = (*self.node_data).clone();
160        let payload = match request.payload.clone() {
161            RequestPayload::Ping => ResponsePayload::Pong,
162            RequestPayload::Store(key, value) => {
163                self.storage.lock().unwrap().insert(key, value);
164                ResponsePayload::Pong
165            }
166            RequestPayload::FindNode(key) => ResponsePayload::Nodes(
167                self.routing_table
168                    .lock()
169                    .unwrap()
170                    .get_closest_nodes(&key, REPLICATION_PARAM),
171            ),

```

(It can be done because we are able to calculate the inter-nodes distance).

Find_value :

Same as find_node, a request of type FindValue is sent:

```

276    /// Sends a `FIND_VALUE` RPC.
277    fn rpc_find_value(&mut self, dest: &NodeData, key: &Key) -> Option<Response> {
278        self.send_request(dest, RequestPayload::FindValue(*key))
279    }

```

Then, when a FindValue request is received, 2 options:

- We can return the value
- We return the closest nodes able to contain the value

```

172        RequestPayload::FindValue(key) => {
173            if let Some(value) = self.storage.lock().unwrap().get(&key) {
174                ResponsePayload::Value(value.clone())
175            } else {
176                ResponsePayload::Nodes(
177                    self.routing_table
178                        .lock()
179                        .unwrap()
180                        .get_closest_nodes(&key, REPLICATION_PARAM),
181                )
182            }

```

Ping :

We send a request to the target of type ping:

```
261     /// Sends a `PING` RPC.
262     fn rpc_ping(&mut self, dest: &NodeData) -> Option<Response> {
263         self.send_request(dest, RequestPayload::Ping)
264     }
```

The target answers PONG:

```
161         RequestPayload::Ping => ResponsePayload::Pong,
```

Join :

The join procedure is done directly when creating a new node. The entry point is the bootstrap parameter:

```
33     /// Constructs a new `Node` on a specific ip and port, and bootstraps the node with an existing
34     /// node if `bootstrap` is not `None`.
35     pub fn new(ip: &str, port: &str, bootstrap: Option<NodeData>) -> Self {
36         let addr = format!("{}", ip, port);
37         let socket = UdpSocket::bind(addr).expect("Error: could not bind to address.");
38         let node_data = Arc::new(NodeData {
39             addr: socket.local_addr().unwrap().to_string(),
40             id: Key::rand(),
41         });
42         let mut routing_table = RoutingTable::new(Arc::clone(&node_data));
43         let (message_tx, message_rx) = channel();
44         let protocol = Protocol::new(socket, message_tx);
45
46         // directly use update_node as update_routing_table is async
47         if let Some(bootstrap_data) = bootstrap {
48             routing_table.update_node(bootstrap_data);
49         }
50
51         let mut ret = Node {
52             node_data,
53             routing_table: Arc::new(Mutex::new(routing_table)),
54             storage: Arc::new(Mutex::new(Storage::new())),
55             pending_requests: Arc::new(Mutex::new(HashMap::new())),
56             protocol: Arc::new(protocol),
57             is_active: Arc::new(AtomicBool::new(true)),
58         };
59
60         ret.start_message_handler(message_rx);
61         ret.start_bucket_refresher();
62         ret.bootstrap_routing_table();
63         ret
64     }
```

Anne Honymé

Line 48, the entry point is added to the bucket table.

Then let's have a look to the *bootstrap_routing_table()* function:

```
107  /// Bootstraps the routing table using an existing node. The node first looks up its id to
108  /// identify the closest nodes to it. Then it refreshes all routing buckets by looking up a
109  /// random key in the buckets' range.
110  fn bootstrap_routing_table(&mut self) {
111      let target_key = self.node_data.id;
112      self.lookup_nodes(&target_key, true);
113
114      let bucket_size = { self.routing_table.lock().unwrap().size() };
115
116      for i in 0..bucket_size {
117          self.lookup_nodes(&Key::rand_in_range(i), true);
118      }
119  }
```

First of all, the node will try to find itself to fill its bucket table. Then it will search for random nodes in order to fulfill the table.

This part could be improved: it could be done concurrently in order to avoid losing time.

Leave:

A **kill** message is sent to the node itself:

```
486  /// Kills the current node and all active threads.
487  pub fn kill(&self) {
488      self.protocol.send_message(&Message::Kill, &self.node_data);
489  }
```

When received, the corresponding node will be marked as inactive by itself, the other nodes are not informed.

```
71      match request {
72          Message::Request(request) => node.handle_request(&request),
73          Message::Response(response) => node.handle_response(&response),
74          Message::Kill => {
75              node.is_active.store(false, Ordering::Release);
76              info!("{}", - Killed message handler", node.node_data.addr);
77              break;
78          }
79      }
80  }
```

This part could be improved, in order to inform the other nodes that we are leaving.

Anne Honymé

Benchmarks :

It looks like the project isn't maintained anymore, the last commit has been pushed 3 years ago. The merge request hasn't been considered.

There are 8 files, sliced into structures and enum (Rust is a functional paradigm language).

There is a total of 1245 lines of an excellent Rust code, including the unit tests; which is quite small.

The nodes communicate through UDP, and the messages are serialized using **serde crate**, in bincode format (<https://docs.rs/bincode/latest/bincode/>), as you can see in the file **src/protocol.rs**.

Criteria for Software Self-Assessment :

kademlia-dht-rs: Family=vehicle; Audience=partners; Evolution=nofuture; Duration=2;
Contribution=none; Url=<https://gitlab.com/jeffrey-xiao/kademlia-dht-rs>

It's a library aiming to create a small implementation of the kademlia protocol in Rust, for educational purposes.